



BIOMOD : Tutorial

Biomod Team :

Wilfried Thuiller, Damien Georges,
Robin Engler, Bruno Lafourcade

February 13, 2012

Contents

1	Before Starting...	3
1.1	Installation and dependencies	3
1.2	General advice	4
1.3	Biomod Contents	5
1.3.1	Biomod Functions	5
1.3.2	Biomod dataset	7
1.3.3	Plotting the data	8
2	Initialisation of Biomod	12
3	Settings in Models()	14
3.1	Calibration and evaluation procedure	14
3.2	Pseudo-absences	18
3.3	Weights	21
4	Running the models	22
4.1	Application of Models()	22
4.2	Going further	24
5	Analysing the outputs	25
5.1	Objects in the workspace	25
5.1.1	Evaluation of the predictive performance	26
5.1.2	Evaluation of the importance of each variable	30
5.1.3	PA data generated	35
5.2	Objects stored on the hard drive : The Models	37
5.2.1	Response Curves	47
5.3	Objects stored on the hard drive : The Predictions	49
5.3.1	Transforming the predictions on the original dataset	57
5.3.2	Identifying the best model	59
6	Uncertainty analysis	64
6.1	Models' projection	64
6.2	Ensemble Forecasting	67
7	Distributions Changes	73
7.1	Species Range Change	73
7.2	Species Turnover	81
7.3	Probability Density Function	84
7.3.1	An example with repetitions	88

1 Before Starting...

In order to facilitate the learning of BIOMOD, a tutorial is provided here with artificial data. It is recommended that the user follows each step and run the models on these artificial datasets, or at least in parallel with runs on its own data. The completion of the tutorial should bring sufficient answers as for the usage of BIOMOD on other datasets.

1.1 Installation and dependencies

To run BIOMOD, please use the latest version of R. A certain number of libraries are also required (rpart, MASS, gbm, gam, nnet, mda, randomForest, Hmisc, plyr) and are also to be downloaded from Rcran before attempting to run BIOMOD.

Note that BIOMOD now enables to build projections directly on rasters. This recent innovation requires several more packages, even if you will not be using rasters with your own work. These are : foreign, sp, rgdal, raster, maptools, some of which are on Rcran and others on the R-forge website.

R code

```
# get list of packages already installed on my computer
myPackages <- .packages(all = TRUE)
biomodDependencies <- c('rpart', 'MASS', 'gbm', 'gam', 'nnet', 'mda',
                        'randomForest', 'Hmisc', 'plyr', 'foreign',
                        'sp', 'rgdal', 'raster', 'maptools')

# compare my packages with those required by BIOMOD
missingPackages <- biomodDependencies[!(biomodDependencies
                                         %in% myPackages)]

# uncomment install.packages() and choose a mirror to install
# missing components
if(length(missingPackages) > 0){
  cat(toString(missingPackages), ' packages have to be installed\n')
  #install.packages(missingPackages, dependencies=T)
} else{ cat('All dependences are installed\n')}
```

R code

```
All dependences are installed
```

BIOMOD is a developing R package that is to be downloaded from the project web page (https://r-forge.r-project.org/R/?group_id=302) and install manually or typing in R :

R code

```
# is biomod already installed
if(!('BIOMOD' %in% myPackages)){
  install.packages("BIOMOD", repos="http://R-Forge.R-project.org")
} else { cat('BIOMOD is already installed')}
```

BIOMOD is already installed R code

It is advised to check relatively frequently for updates.

```
# update biomod if necessary
update.packages("BIOMOD", repos="http://R-Forge.R-project.org")
```

1.2 General advise

The recommended procedure is to first create a working directory, for example called BIOMOD. Then, create a new folder where to store the datasets, run the models and save the outputs and results. In our examples, we will create and use the directory called Biomod runs. It is from this folder that the files will be read and written. You need to put a copy of your datasets in order to be able to open them once the working directory in R is set to this workspace.

If you want to pause and continue work on this tutorial (or your own project) later. Just save your session. You will get back all your working space just loading the created file.

```
# save all the working space
save.image("BiomodTutorial.RData")
# free the working space
rm(list=ls())
# and get it back
load("BiomodTutorial.RData")
```

Do keep in mind that some information is kept in the file that has just been generated but that a lot of our work is also stored in the directories that have been created by BIOMOD. Both will be needed for carrying on the next steps.

1.3 Biomod Contents

1.3.1 Biomod Functions

The first thing to do is to load the BIOMOD package. It will load all the functions required to run BIOMOD as well as the examples files to be used in this practical.

R code

```
# load BIOMOD package
library(BIOMOD)
```

R code

```
Loaded gbm 1.6-3.1
```

To access all BIOMOD functions :

R code

```
# listing of BIOMOD functions
help(package='BIOMOD')
```

As for any function, you can access help files :

R code

```
?response.plot
```

```
response.plot (BIOMOD) R Documentation
```

Analysis of the response curves of a model within Biomod

Description

Adaptation of the Evaluation Strip proposed by Elith et al.(2005). This function enables to plot the response curves of a model independently of the algorithm used for building the model. It therefore permits a direct comparisons of models built using different statistical approaches on the same data.

Usage

```
response.plot(model, Data, show.variables=seq(1:ncol(Data)), save.file="no", name="response_curve", ImageSize=480)
```

Arguments

model the model for which you want the response curves to be plotted. Compatible with GAM, GBM, GLM, ANN, CTA, RF, FDA and MARS.

Data the variables for which you want the response curves to be plotted. A data frame is wanted with one column per variable. They have to have the same names as the ones used to calibrate the model.

show.variables give in the column numbers of 'Data' for selecting the variables that are wanted for plotting

save.file can be set to "pdf", "jpeg" or "tiff" to save the plot. Pdf options can be changed by setting the default values of pdf.options().

name the name of the file produced if save.file is different to "no" (extensions are already included)

ImageSize the size of the image in pixels if save.file is different to "no". Affects "jpeg" and "tiff" outputs only. Default if 480 pixels which is the R default.

Details

For building such response curves, n-1 variables are set to their median value and only the one of interest is varying across its whole range. The variations observed and the curve thus obtained shows the sensibility of the model to that specific variable. This method does therefore not account for interactions between variables.

Author(s)

Wilfried Thuiller, Bruno Lafourcade

References

Elith, J., Ferrier, S., Huettmann, F., & Leathwick, J. R. 2005 The evaluation strip: A new and robust method for plotting predicted responses from species distribution models. *Ecological Modelling* 186, 280-289.

See Also

[Models](#)

You can also open the Biomod pdfs directly from R :

```
#to open the old but detailed version  
Biomod.Manual()  
#to open one of the latest versions : several pdf files  
Biomod.Manual("Biomod_Presentation_Manual")
```

BIOMOD is composed of a series of functions that enables to do our species modelling :

- **Running BIOMOD**

- Initial.State
- Models
- Projection
- Ensemble.Forecasting

- **Further BIOMOD steps**

- CurrentPred
- PredictionBestModel
- ProjectionBestModel
- Biomod.Turnover
- Biomod.RangeSize
- Migration

- **Plotting functions**

- level.plot
- multiple.plot
- response.plot

- **Other functions**

- ProbDensFunc : *calculates density probabilities*
- pseudo.abs : *generating pseudo-absences*
- BiomodManual : *opens the pdf manual and practicals from R*

We will mainly focus here on the *Models* function as it contains all the options for calibrating and evaluating the models and look at how it can lead to significant variability in prediction making. This function runs the models and evaluation technics presented in the Presentation Manual of BIOMOD (see *Biomod.Manual('Presentation')*).

1.3.2 Biomod dataset

We need to import the species and the environmental data for our modelling. In our example the same file holds the two datasets.

```

R code
# Loading the example datasets
# For practical reasons, species and environment datasets
# are stored together
data(Sp.Env)
head(Sp.Env)

```

	Idw	X	Y	Var1	Var2	Var3	Var4	Var5	Var6
1	73	-9.288	38.62	0.6683	4296	770.1	39.33	295.1	16.74
2	74	-9.292	39.52	0.7596	4174	928.1	57.32	348.7	16.41
3	75	-9.290	39.07	0.7424	4173	870.3	50.05	330.0	16.41
4	76	-8.715	37.72	0.5543	4264	620.0	24.99	239.1	16.66
5	77	-8.717	37.27	0.5489	4169	622.3	25.16	241.0	16.40
6	78	-8.148	37.72	0.5363	4206	591.8	25.74	222.9	16.49

	Var7	Sp281	Sp290	Sp277	Sp164	Sp163	Sp177	Sp185	Sp191
1	10.87	0	1	0	0	1	0	0	1
2	10.51	0	1	0	0	1	0	0	1
3	10.50	0	0	0	0	1	0	0	1
4	10.93	0	0	0	0	0	0	0	0
5	11.28	0	0	0	0	0	0	0	0
6	10.13	0	0	0	0	0	0	0	0

- **Idw**: An Id to keep track of the row numbers
- **X and Y**: longitude and latitude of our sites (for plots, not needed for the modelling in itself)
- **Var1 to Var7** : Environmental variables (bioclimatic in that case)
- **Sp281 to Sp191**: Presence/absence of 8 species

To avoid, confusion, we will split the dataset into 3 part :

- the points coordinates (*LatLong*)
- the bioclimatic data (*Expl.Var*)
- the species occurrences(*Resp.Var*)

```
R code
#Visualisation of our data (show first six rows)
LatLong <- Sp.Env[,2:3] # coordinates of points
Expl.Var <- Sp.Env[,4:10] # bioclimatic variables
Resp.Var <- Sp.Env[,11:17] # species occurrences
```

BIOMOD does not read the coordinates and does not recognise any geographical information when proceeding the modelling. The user should ensure that all datasets are kept in the same order, i.e. each species information (presence or absence) is correctly associated to the explanatory variables. Any mismatch will not be recognised by BIOMOD and the influence on the different outputs and results will be unnoticeable but real.


To load your own data from a text file, use the *read.table()* function:

```
R code
#Loading from a text file
#My.Data <- read.table("my_data.txt", h=T, sep="\t")
```

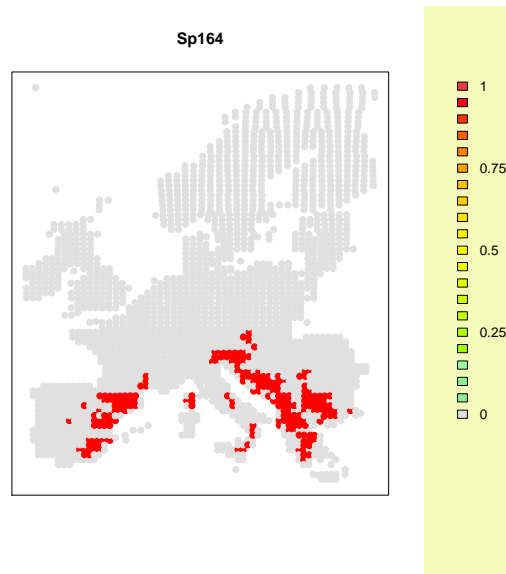
1.3.3 Plotting the data

The *level.plot* function requires two inputs : the vector of values that you want to plot and the coordinates of your data points. It works with any type of data.

```
R code
level.plot(Expl.Var[,4], LatLong[,1:2], title=colnames(Expl.Var)[4])
```

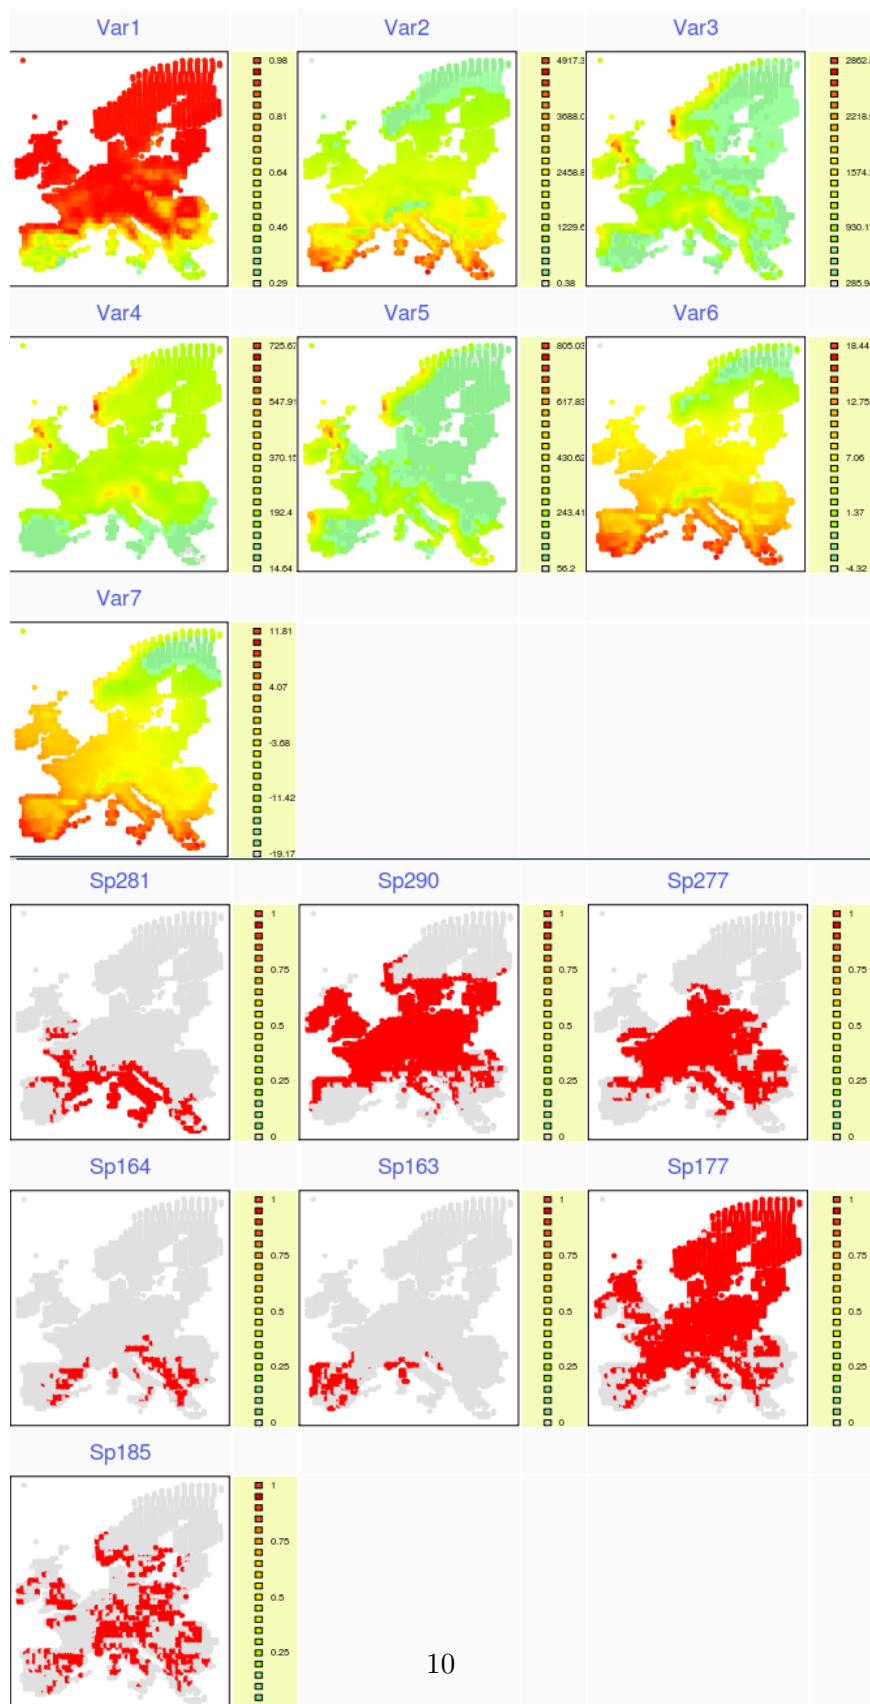


```
R code  
level.plot(Resp.Var[,4], LatLong[,1:2], title=colnames(Resp.Var)[4])
```



Let's take a general view of our data with the `multiple.plot` function :

```
R code  
multiple.plot(Expl.Var, LatLong[,1:2], cex=0.7)  
multiple.plot(Resp.Var, LatLong[,1:2], cex=0.7)
```



You can modify the color gradient by setting the *color.gradient* argument to either *red* (the default), *blue* or *grey*.

2 Initialisation of Biomod

First, we need to set up the dataset in a correct format for BIOMOD by means of the *Initial.State* function. The syntax in the function is the following:

- **Response** : The response variables to model.
- **Explanatory**: The explanatory or independent variables.

Additional arguments (see the *Presentation* pdf for explanation) :

- **IndependentResponse** : Truly independent response variables.
- **IndependentExplanatory** : Truly independent explanatory variables.

These are used to evaluate the predictive accuracy of the models.

We will work on *Sp.Env* dataset, see 1.3.2 to load it correctly or adapt the following code lines to fit with your own data.

So our call looks like :

R code

```
Initial.State(Response = Resp.Var[,1:2], Explanatory = Expl.Var)
```

But we will inform anyway the 2 optional arguments with the same information. The point is to have an example of predictions on our full database as we are going to use pseudo-absences for the purpose of the example (hence BIOMOD will only produce predictions on partial data).

So instead we have :

R code

```
Initial.State(Response = Resp.Var[,1:2],
              Explanatory = Expl.Var,
              IndependentResponse = Resp.Var[,1:2],
              IndependentExplanatory = Expl.Var)

ls()
```

R code

```
[1] "biomodDependencies" "Biomod.material"
[3] "DataBIOMOD"        "DataEvalBIOMOD"
[5] "Expl.Var"          "LatLong"
[7] "missingPackages"   "myPackages"
[9] "Resp.Var"          "Sp.Env"
```

It creates 'DataBIOMOD' our reference database, and DataEvalBIOMOD if you have given independent information. The latter will be used during the testing of the models. Make sure to always keep these datasets unchanged and never delete them.

R code

```
head(DataBIOMOD)
```

	Var1	Var2	Var3	Var4	Var5	Var6	Var7	Sp281	Sp290
1	0.6683	4296	770.1	39.33	295.1	16.74	10.87	0	1
2	0.7596	4174	928.1	57.32	348.7	16.41	10.51	0	1
3	0.7424	4173	870.3	50.05	330.0	16.41	10.50	0	0
4	0.5543	4264	620.0	24.99	239.1	16.66	10.93	0	0
5	0.5489	4169	622.3	25.16	241.0	16.40	11.28	0	0
6	0.5363	4206	591.8	25.74	222.9	16.49	10.13	0	0

DataBIOMOD contains the environmental variables in the first columns, followed by the species occurrences. DataEvalBIOMOD has the same structure but it contains the data for testing the models.

An object called Biomod.material is also produced which contains information that has been extracted from the datasets like the number of variables, the number of species, etc.. Most of the functions will refer to this object to obtain some necessary values, so make sure to keep it unchanged.

R code

```
Biomod.material
```

R code

```
$NbVar
[1] 7

$VarNames
[1] "Var1" "Var2" "Var3" "Var4" "Var5" "Var6" "Var7"

$NbSpecies
[1] 2

$species.names
[1] "Sp281" "Sp290"
```

3 Settings in Models()

The *Models()* function will run the different models available in BIOMOD and described in the *Presentation* manual. There are two main issues to consider : which models to select and what calibration/evaluation procedure to choose. Let's first have a look at the options to be set in the *Models()* function (arguments are presented with their default values):

```
Models(  
  Setting the models to TRUE or FALSE (to run them or not) and their as-  
  sociated options (please refer to the Presentation Manual)  
  GLM=FALSE, TypeGLM="simple", Test="AIC",  
  GBM=FALSE, No.trees= 5000,  
  GAM=FALSE, Spline=3,  
  CTA=FALSE, CV.tree=50,  
  ANN=FALSE, CV.ann=5,  
  SRE=FALSE, quant=0.025,  
  FDA=FALSE,  
  MARS=FALSE,  
  RF=FALSE,
```

The calibration procedure options

```
NbRunEval=1, DataSplit=100,  
NbRepPA=0, strategy="sre", coor=NULL, distance=0, nb.absences=NULL,  
Yweights=NULL,
```

The evaluation procedure options

```
VarImport=0,  
Roc=FALSE, Optimized.Threshold.Roc=FALSE, Kappa=FALSE, TSS=FALSE,  
KeepPredIndependent=FALSE  
)
```

Note that the various models' specific options will directly influence **the inner** calibration procedure of the models, whereas the calibration options below (NbRunEval, DataSplit) determine **the general trend** of the calibration which will be applied to all the models in the same way.

3.1 Calibration and evaluation procedure

A key issue in modelling is the calibration procedure of the models with the constant effort to obtain a reliable estimation of their performance.

Ideally, one should always evaluate the predictive performance of a model using independent data, i.e. data from which the model didn't obtain any information to build itself. this would enable to reliably test its predictive accuracy on a new dataset and certify its efficiency. Unfortunately, this kind of information is rarely accessible in species distribution modelling. An alternative to assess the predictive performance of the models is to split the original data in calibration (training) and evaluation (testing) datasets : one part is used to feed the model, the other, kept aside and therefore new to the model, is used to check the models' efficiency to predict the right value. As a consequence, this method consists of a trade-off between the amount of data used for the construction of the model and the accuracy of the evaluation measure.

This splitting procedure, widely used in the modelling world, nevertheless brings a major issue : the subsequent randomness of the data selection used for calibration and its impact on the modelling quality.

To obtain a reliable way of evaluating the models while not influencing the prediction making by the random splitting of the data, BIOMOD proposes to build a series of models. The above calibration/evaluation procedure is repeated a certain number of times to perform a reliable evaluation as an attempt to free ourselves from the random effect (the mean result is extracted). Then a final model is built without splitting the data, i.e. 100 % of the data available is used, thus using all the information available and not having any random effect in the prediction making.

This method is also a good way of assessing for uncertainty. While many modellers are satisfied with running only their models once, we propose to build a large number of models to measure the sensitivity of the models to the initial conditions (the input data given). Each model built is kept and can be used to later render projections.

The combination of the two arguments below will determine in which way the models will be built and tested.

- **NbRunEval**: number of random data splitting procedure for creating calibration and evaluation datasets ; a model will be built from each one of them. If set to zero, only the final 100 % model is built.

- **DataSplit**: the ratio used for splitting the original database in calibration and evaluation subsets (value to give is the % awarded for calibration). A 70/30 % partitioning is recommended as commonly used (Arajo, et al. 2005b, Guisan and Thuiller 2005).

pros : It gives a more robust estimate of the predictive performance of each selected model and it also provides an assessment of the sensitivity of the model to the initial conditions, i.e. to the species distribution data.

cons : it lengthens the modelling time needed to build the models (it can be an exceeding amount of time if not done carefully).

main interest : adds variability in the predictions when several runs are made due to the random effect of selecting the data, i.e. each model is not build using the same information, representing the sensibility of the models on the input data.

Example with the fda and species Sp281

Here is an example of the effect of randomness in the prediction making.

R code

```
#to call our dataset
# library(BIOMOD)
# data(Sp.Env)

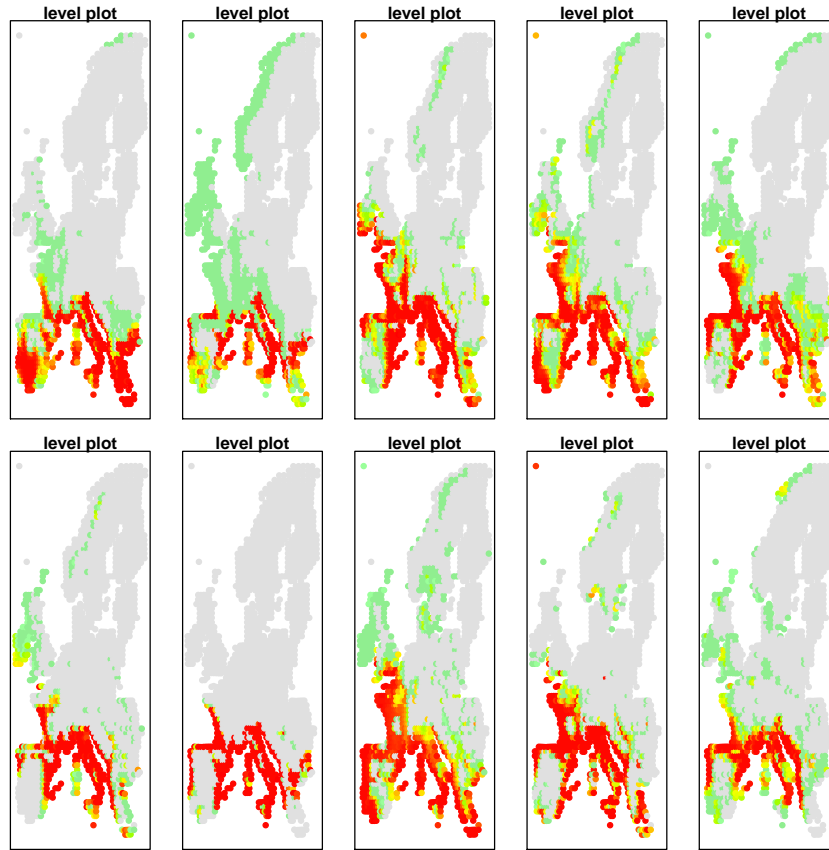
store <- matrix(nr=nrow(Resp.Var), nc=0)
for(i in 1:10){
  rand <- sample(nrow(Resp.Var), 100)
  model <- fda("Sp281 ~Var1 + Var2 + Var3 + Var4 + Var5 + Var6 + Var7",
              data=cbind(Expl.Var[rand,],Resp.Var[rand,]), method=mars)
  store <- cbind(store, predict(model, Sp.Env[,4:10], type="post")[,2])
}
```

R code

```
for(i in 1:10){
  x11()
  par(mar=c(1,1,1,1))
  level.plot(store[,i], LatLong)
}
```

R code

```
par(mfrow=c(2,5))
par(mar=c(1,1,1,1))
for(i in 1:10) level.plot(store[,i], LatLong, show.scale=F, cex=0.85)
```



This is the same model (FDA) and the same datasets used, only the initial calibration data is changing. The impact on the geographical patterns can clearly be seen.

3.2 Pseudo-absences

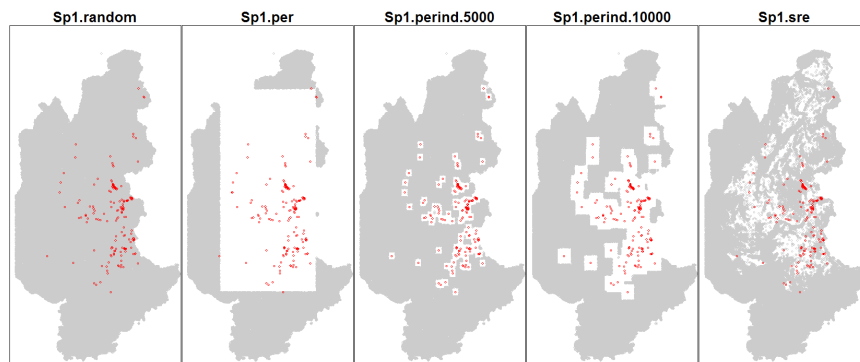
All the models in BIOMOD need information about presences and absences for being able to determine the suitable conditions for a given species. Some datasets, however, do not contain absences but only presences and the construction of virtual absences is therefore needed. This is, for example, the case of bird datasets where determining an absence can be rather tricky. The assumed absences are called pseudo-absences for there is no field verification of this generated information.

These pseudo-absences are created by considering any point where the species was not recorded and where the environmental conditions are known to cause potential absence. Feeding the models with exceeding numbers of absences can significantly disturb the ability of models to discriminate meaningful relationships between climate and species distributions. Moreover, running

models on such heavy databases is incredibly time consuming.

In addition, some of the chosen absences might unfortunately represent true presences (this is particularly likely in the case of incomplete samples) and therefore the pseudo-absence data gives false information for the estimation of the species-climate relationship. Hence, we propose various strategies that seek to remove the spurious effects of using poorly selected pseudo-absences before running the models.

Example of the 4 available strategies in the region of the French Alps for *Larix decidua miller*. The presences are in red and the pseudo-absences selected by each strategy are in grey.



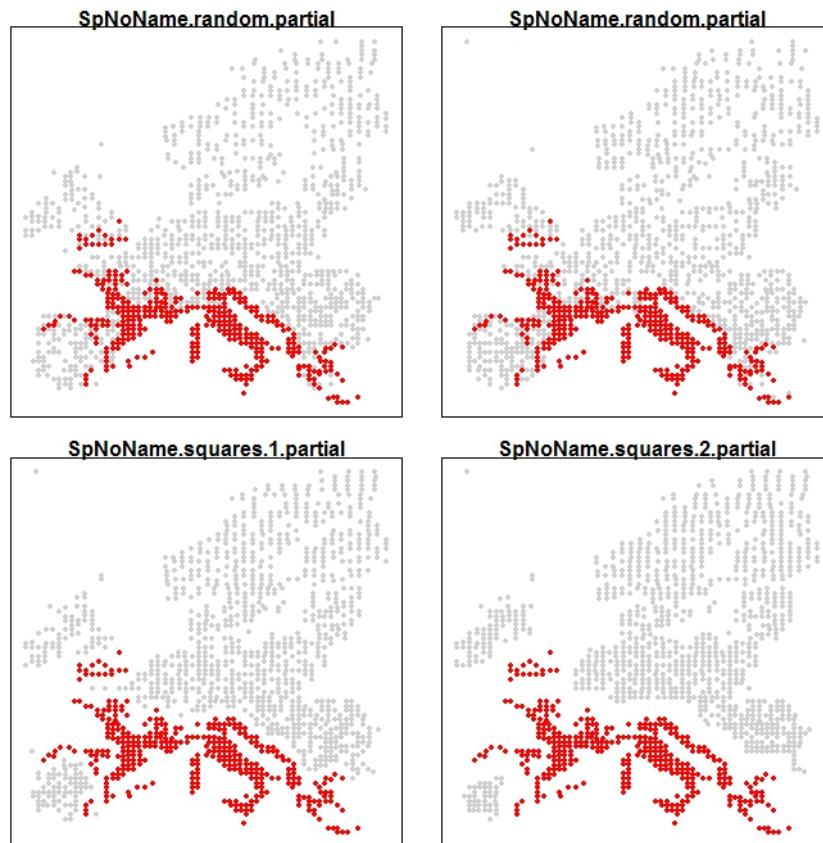
A few examples of what datasets would be created in our case :

R code

```

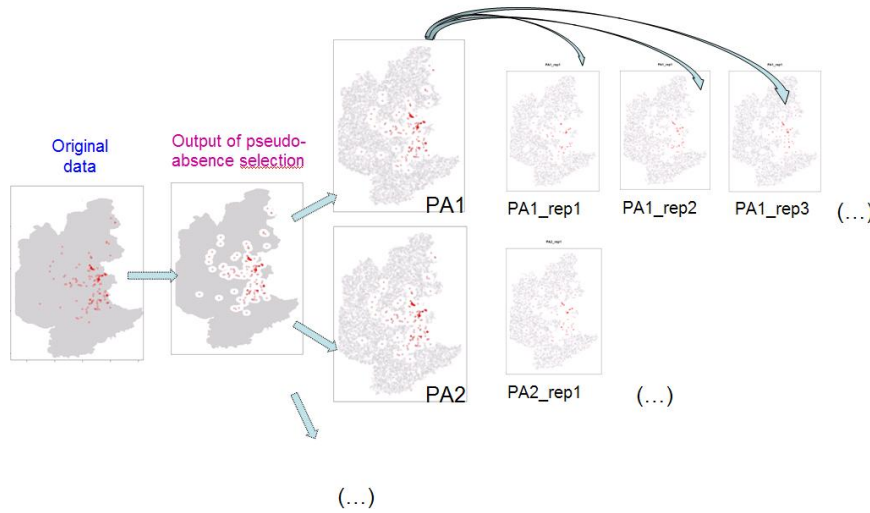
par(mfrow=c(2,2), mar=c(1,1,1,1))
x <- pseudo.abs(LatLong, Resp.Var[, "Sp281"], strategy='random',
  plot=T, nb.points=1000)
x <- pseudo.abs(LatLong, Resp.Var[, "Sp281"], strategy='random',
  plot=T, nb.points=1000)
x <- pseudo.abs(LatLong, Resp.Var[, "Sp281"], strategy='squares',
  plot=T, nb.points=1000, distance=1)
x <- pseudo.abs(LatLong, Resp.Var[, "Sp281"], strategy='squares',
  plot=T, nb.points=1000, distance=2)

```



In *Models()*, you can choose to run pseudo-absences selections with the argument *NbRepPA*.

This argument is to be correlated with the usage of repetitions for the calibration : once the pool of potential pseudo-absences has been defined by the strategy selected, a user-defined number (*Nb.absences* argument) is randomly selected from this pool. We therefore have a random effect in the calibration process coming from the creation of pseudo-absences for our data. The *NbRepPA* argument will define a number of repetitions for randomly withdrawing absences to constitute the calibration datasets. Do consider that the total number of repetitions will be a multiplication of the two repetition arguments



3.3 Weights

The *Yweights* argument enables the user to set extra information for the response variables (a matrix with N columns for the N species). This is similar to an index of detectability for each site, which allows users to give stronger weights to more reliable presences or absences. It can be scaled up and put as a weight in the modeling process. For more information, see how *weights* is working in R.

4 Running the models

4.1 Application of Models()

We can now run the different models on our species. It takes only a few moments for each model to run. All the selected models (= TRUE) will run for each species. Here we will have 9(models selected)*4(3 repetitions + final model)*2(PA repetitions) which makes 72 models per species, it will thus take several minutes.

Please, be aware that the *NbRunEval* and *NbRepPA* arguments can considerably enlarge your calculation time by multiplying the number of runs to be made for each species. Do not enter excessively high values for these two arguments **unless** you have sufficient patience and/or reasonable calculation power.

```

Models(GLM = T, TypeGLM = "poly", Test = "AIC",
       GBM = T, No.trees = 2000,
       GAM = T, Spline = 3, CTA = T, CV.tree = 50,
       ANN = T, CV.ann = 2,
       SRE = T, quant=0.025,
       FDA = T,
       MARS = T,
       RF = T,
       NbRunEval = 3, DataSplit = 80, Yweights=NULL,
       Roc = T, Optimized.Threshold.Roc = T, Kappa = T, TSS=T,
       KeepPredIndependent = T, VarImport=5,
       NbRepPA=2, strategy="circles", coor=LatLong,
       distance=2, nb.absences=1000)

```

For the purpose of the example (even though the data does not ask for it) we used 2 pseudo-absences (PA) runs. Note that there has only been one PA run for Sp290 because too little absences were available compared to the ones wanted. The nb.absences argument was set to 1000, but:

```

#the number of data selected by the pseudo-absences procedure
length(Biomod.PA.data$Sp290)

```

```

[1] 1773

```

R code

```
#the number of presences for Sp290
sum(Sp.Env[, "Sp290"])
```

R code

```
[1] 1350
```

R code

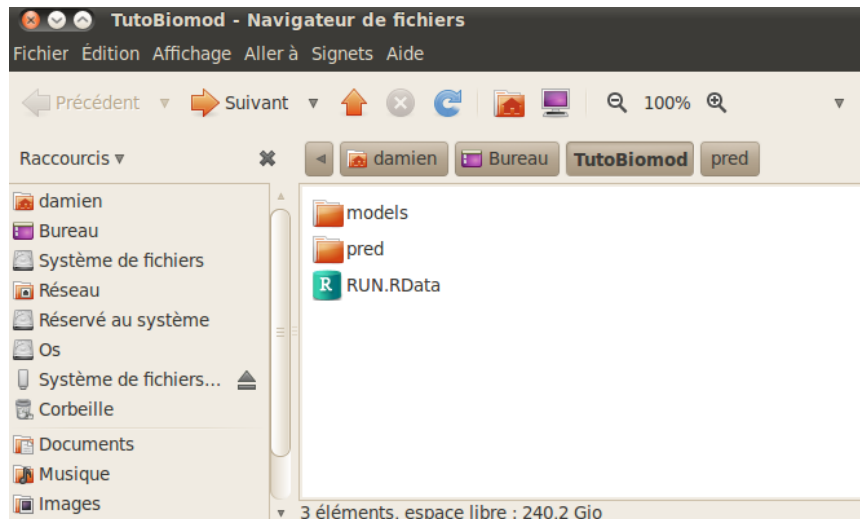
```
#Hence, the number of absences available for calibration
length(Biomod.PA.data$Sp290) - sum(Sp.Env[, "Sp290"])
```

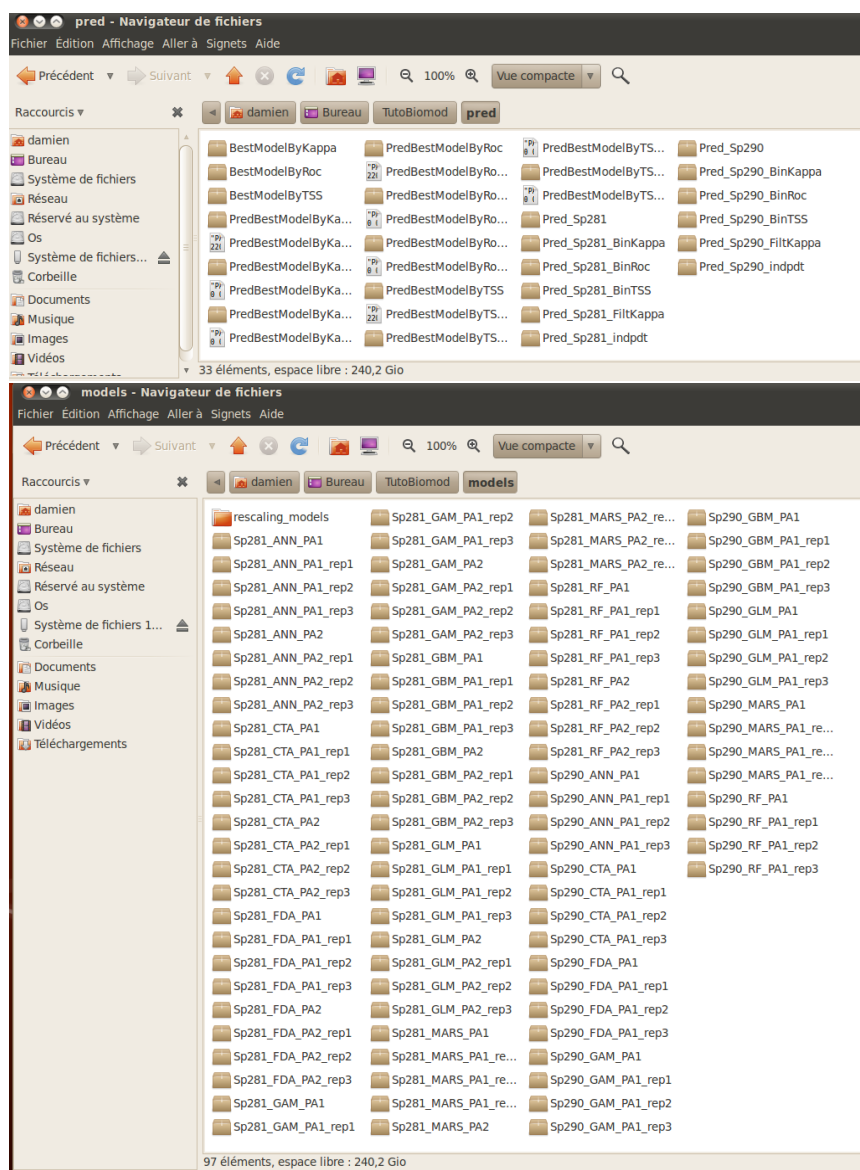
R code

```
[1] 423
```

Too little absences are available. In this case, a single pseudo-absences run is made using all the absences available.

In the latter version of BIOMOD, the results are stored outside R's workspace to counter the memory storage limitations of the software. While running BIOMOD, you will realise that additional folders will be created. A series of objects have been produced in the workspace and also on the harddrive of your computer. Your working folder should now look like this.





4.2 Going further

For those which are interesting in how each model is computed in BIOMOD, you can have a look on the last BIOMOD summer school 'Methods' practical (*W. Thuiller*). An archived file containing script 'Methods.r' and data required may have been send you with this tutorial.

5 Analysing the outputs

5.1 Objects in the workspace

There are now various objects stored in the workspace. First, we can have a look at what is present in our R session and check what has been produced by the *Models()* function.

<code>ls()</code>	R code
[1] "BestModelByRoc"	"BestModelByTSS"
[3] "biomodDependencies"	"Biomod.material"
[5] "Biomod.PA.data"	"Biomod.PA.sample"
[7] "DataBIOMOD"	"DataEvalBIOMOD"
[9] "data.used"	"Evaluation.results.Kappa"
[11] "Evaluation.results.Roc"	"Evaluation.results.TSS"
[13] "Expl.Var"	"Expl.Var2"
[15] "Expl.Var3"	"Future1"
[17] "GBM.list"	"GBM.perf"
[19] "i"	"isnulllyweights"
[21] "LatLong"	"missingPackages"
[23] "model"	"myPackages"
[25] "obj"	"our.lines"
[27] "Pred"	"Pred2"
[29] "Pred3"	"PredBestModelByKappa"
[31] "Pred_Sp281"	"Pred_Sp290"
[33] "Pred_Sp290_BinKappa"	"Pred_Sp290_FiltKappa"
[35] "Pred_Sp290_indpdt"	"rand"
[37] "Resp.Var"	"Sp290_GLM_PA1"
[39] "Sp290_RF_PA1"	"Sp.Env"
[41] "store"	"VarImportance"

So, we have the outputs generated by *Initial.State* and the original datasets :

- Sp.Env
- LatLong
- Expl.Var
- Resp.Var
- DataBIOMOD
- Biomod.material

We also have the objects produced by the *Models()* function in the workspace (additional objects are stored on the hard disk). These are :

- Evaluation.results.Roc
- Evaluation.results.Kappa
- Evaluation.results.TSS
- VarImportance.

And we get the following if NbRepPA is higher than 0 :

- Biomod.PA.data
- Biomod.PA.sample
- SpNoName.circles.2 (or something close)

5.1.1 Evaluation of the predictive performance

There are three available techniques for making an assessment of a model's performance. A summary table of the type "Evaluation.results.method" are produced containing the predictive performance of each model which is convenient for making comparisons across methods and taxa.

R code

```
#Here we only display the info for the first species modelled
Evaluation.results.Kappa[1:8]
```

R code

```
$Sp281_PA1
Cross.validation indepth.data total.score Cutoff
ANN          0.879          0.625          0.9075 438.0
CTA           0.883          0.614          0.9579 630.0
GAM           0.855          0.674          0.8829 629.4
GBM           0.901          0.646          0.9148 592.8
GLM           0.894          0.65           0.8490 699.3
MARS          0.926          0.676          0.9200 429.6
FDA           0.880          0.642          0.9170 105.8
RF            0.930          0.763          1.0000 340.0
SRE           0.658          0.394          0.6675 10.0

Sensitivity Specificity
ANN          96.17          96.2
CTA           98.98          98.0
GAM           94.13          95.6
```

GBM	97.19	96.2
GLM	89.03	95.8
MARS	94.13	97.8
FDA	94.90	97.3
RF	100.00	100.0
SRE	83.42	86.8

\$Sp281_PA1_rep1

	Cross.validation	indepdt.data	total.score	Cutoff
ANN	0.841	none	0.8929	171.3
CTA	0.858	none	0.8748	718.5
GAM	0.839	none	0.8561	619.4
GBM	0.896	none	0.9156	654.3
GLM	0.841	none	0.8408	769.2
MARS	0.956	none	0.9281	599.4
FDA	0.853	none	0.8852	109.8
RF	0.930	none	0.9858	410.0
SRE	0.669	none	0.6415	10.0

	Sensitivity	Specificity
ANN	99.23	94.0
CTA	98.72	93.1
GAM	93.37	94.3
GBM	95.41	97.0
GLM	85.71	96.8
MARS	92.86	98.8
FDA	92.60	96.4
RF	99.49	99.4
SRE	83.93	84.7

\$Sp281_PA1_rep2

	Cross.validation	indepdt.data	total.score	Cutoff
ANN	0.930	none	0.9265	431.6
CTA	0.911	none	0.9403	630.0
GAM	0.855	none	0.8848	609.4
GBM	0.895	none	0.9152	639.2
GLM	0.929	none	0.9278	659.3
MARS	0.900	none	0.9230	239.8
FDA	0.876	none	0.9124	228.6
RF	0.921	none	0.9841	330.0
SRE	0.633	none	0.6936	10.0

	Sensitivity	Specificity
ANN	96.94	97.0
CTA	97.45	97.6
GAM	94.39	95.6
GBM	94.64	97.3
GLM	96.17	97.4
MARS	96.68	96.9
FDA	92.60	98.0
RF	99.23	99.4

SRE 81.63 89.5

\$Sp281_PA1_rep3

	Cross.validation	indepdt.data	total.score	Cutoff
ANN	0.866	none	0.8951	395.2
CTA	0.879	none	0.9290	340.0
GAM	0.870	none	0.8946	569.4
GBM	0.913	none	0.9124	612.9
GLM	0.913	none	0.9328	669.3
MARS	0.921	none	0.9243	629.4
FDA	0.911	none	0.9248	302.9
RF	0.938	none	0.9876	420.0
SRE	0.672	none	0.6752	10.0

	Sensitivity	Specificity
ANN	94.90	96.0
CTA	98.98	96.3
GAM	96.68	95.2
GBM	95.92	96.6
GLM	95.92	97.8
MARS	92.09	98.9
FDA	93.37	98.4
RF	99.49	99.5
SRE	84.44	86.8

\$Sp281_PA2

	Cross.validation	indepdt.data	total.score	Cutoff
ANN	0.868	0.639	0.9493	293.20
CTA	0.868	0.622	0.9309	210.00
GAM	0.848	0.678	0.8859	749.25
GBM	0.903	0.641	0.9235	577.40
GLM	0.763	0.652	0.8437	749.25
MARS	0.921	0.686	0.9345	359.64
FDA	0.916	0.653	0.9118	72.12
RF	0.941	0.767	1.0000	390.00
SRE	0.669	0.394	0.6546	10.00

	Sensitivity	Specificity
ANN	98.72	97.6
CTA	99.49	96.2
GAM	88.78	98.1
GBM	97.96	96.4
GLM	86.99	96.4
MARS	95.66	98.0
FDA	94.64	97.1
RF	100.00	100.0
SRE	83.42	85.9

\$Sp281_PA2_rep1

	Cross.validation	indepdt.data	total.score	Cutoff
ANN	0.854	none	0.9063	402.0

CTA	0.860	none	0.9096	630.0
GAM	0.829	none	0.8566	558.9
GBM	0.885	none	0.9089	604.1
GLM	0.758	none	0.7735	625.0
MARS	0.918	none	0.9258	659.3
FDA	0.909	none	0.9226	389.7
RF	0.937	none	0.9876	450.0
SRE	0.714	none	0.6821	10.0

Sensitivity Specificity

ANN	96.94	95.8
CTA	96.94	96.0
GAM	95.92	93.2
GBM	95.66	96.5
GLM	90.05	90.6
MARS	91.58	99.2
FDA	92.35	98.7
RF	98.98	99.7
SRE	82.65	88.2

\$Sp281_PA2_rep2

Cross.validation indepdt.data total.score Cutoff

ANN	0.865	none	0.8919	591.9
CTA	0.851	none	0.9228	660.0
GAM	0.807	none	0.8390	608.8
GBM	0.886	none	0.9061	657.3
GLM	0.724	none	0.7552	688.6
MARS	0.881	none	0.9268	399.6
FDA	0.884	none	0.9232	179.9
RF	0.913	none	0.9805	380.0
SRE	0.646	none	0.6505	10.0

Sensitivity Specificity

ANN	97.70	94.6
CTA	96.43	97.0
GAM	92.35	93.7
GBM	93.62	97.2
GLM	85.71	91.5
MARS	93.88	98.3
FDA	93.62	98.2
RF	99.23	99.2
SRE	84.18	85.2

\$Sp281_PA2_rep3

Cross.validation indepdt.data total.score Cutoff

ANN	0.884	none	0.9259	404.0
CTA	0.893	none	0.8865	726.8
GAM	0.908	none	0.8517	598.8
GBM	0.938	none	0.9148	603.8
GLM	0.807	none	0.7467	706.4
MARS	0.965	none	0.9242	329.7

<i>FDA</i>	<i>0.956</i>	<i>none</i>	<i>0.9219</i>	<i>118.2</i>
<i>RF</i>	<i>0.973</i>	<i>none</i>	<i>0.9947</i>	<i>490.0</i>
<i>SRE</i>	<i>0.646</i>	<i>none</i>	<i>0.6145</i>	<i>10.0</i>
	<i>Sensitivity</i>	<i>Specificity</i>		
<i>ANN</i>	<i>95.66</i>	<i>97.5</i>		
<i>CTA</i>	<i>96.94</i>	<i>94.6</i>		
<i>GAM</i>	<i>93.88</i>	<i>93.8</i>		
<i>GBM</i>	<i>97.19</i>	<i>96.2</i>		
<i>GLM</i>	<i>84.18</i>	<i>91.7</i>		
<i>MARS</i>	<i>95.66</i>	<i>97.4</i>		
<i>FDA</i>	<i>94.39</i>	<i>97.8</i>		
<i>RF</i>	<i>99.49</i>	<i>99.9</i>		
<i>SRE</i>	<i>84.18</i>	<i>82.6</i>		

You can explore and see that the PA2 runs for Sp290 are empty matrices. That's because there has only been 1 PA run for that species.

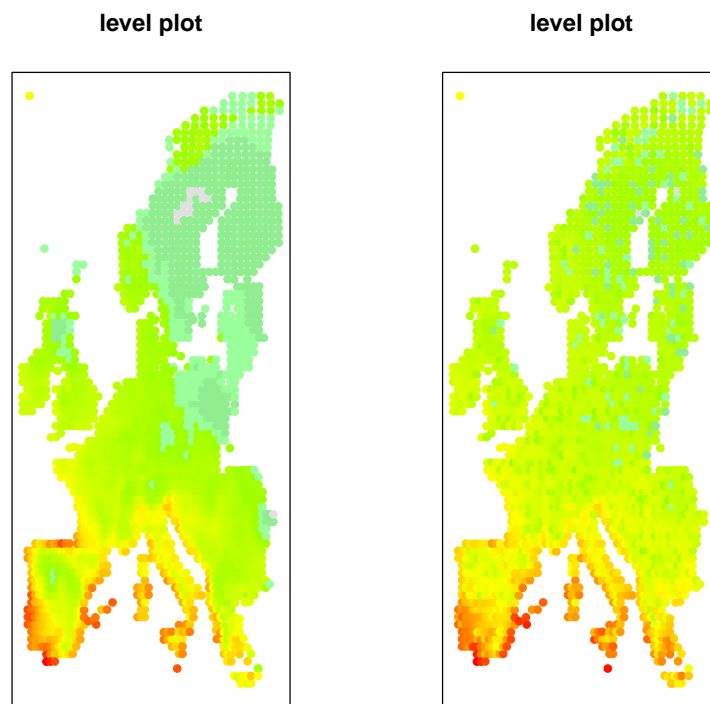
5.1.2 Evaluation of the importance of each variable

It is always difficult to compare predictions from different models as they do not rely on the same algorithms, techniques and assumptions about the expected relationship between the response and the variables, i.e. the species distributions and the environment. With a permutation procedure, BIOMOD proposes another way to examine the importance of the variables in the models. We extract a measure of relative importance of each variable that is independent of the model. Note that the importance of the variables is only calculated for the final model.

Procedure: once the models are trained (i.e. calibrated), a standard prediction is made. Then, one of the variables is randomized and a new prediction is made. The correlation score between that new prediction and the standard prediction is calculated and is considered to give an estimation of the variable importance in the model :

```
R code
model <- glm(Sp281 ~ Var1 + Var2 + Var3 + Var4 + Var5 + Var6 + Var7, data=Sp.Env)
Pred <- predict(model, Expl.Var, type="response")
```


```
R code
Expl.Var2 <- Expl.Var
Expl.Var2[, 'Var1'] <- sample(Expl.Var[, 'Var1'])
Pred2 <- predict(model, Expl.Var2, type="response")
par(mfrow=c(1,2))
level.plot(Pred, LatLong, show.scale=F, cex=0.8)
level.plot(Pred2, LatLong, show.scale=F, cex=0.8)
```



```
cor(Pred, Pred2)
```

```
[1] 0.911
```

```
plot(Pred, Pred2)
```




A good correlation score between the two predictions, i.e. they only slightly differ, shows that the randomized variable has little influence on the prediction making and is considered not important for the model in its prediction.

R code

```
Expl.Var3 <- Expl.Var  
Expl.Var3[, 'Var7'] <- sample(Expl.Var[, 'Var7'])  
Pred3 <- predict(model, Expl.Var3, type="response")  
plot(Pred, Pred3)  
cor(Pred, Pred3)
```

R code

```
[1] -0.268
```



In contrary, a low correlation means a significant difference in the prediction making, showing an importance of that variable for the model.

NOTE : in the *VarImportance* output, the values given correspond to 1 minus the correlation score. High values will therefore reveal a high importance of the variable whereas a value close to 0 will reveal no importance.

Score of variable 1 (Pred2) : $1 - \text{cor}(\text{Pred}, \text{Pred2}) = 0.09$ meaning low influence

Score of variable 2 (Pred3) : $1 - \text{cor}(\text{Pred}, \text{Pred3}) = 1.27$ meaning high influence

This step is repeated n times for each variable independently and the means are kept for each variable.

NOTE : The obtained correlation can be negative. We consider these cases to represent an even bigger influence of the permuted variable on

the prediction than with a correlation of 0. The variable importance estimation will therefore still be given as 1 minus the correlation score and, as a consequence, turn into values higher than 1. These cases are not so rare.

Running the *Models* function will produce an object called "VarImportance" (only if VarImp was put higher than 0 in the function call). The results are stored individually per species and per model. Let's look at the results we have :

Values should be considered independently for each model. For instance, the SRE shows a generally low value for all the variable when the ANN is generally high. The goal is nevertheless to identify which variable is of the most importance. A good example with the GLM for Sp281, only 2 variables seem to have a significance in the predictions.

Note also that this technic only accounts for the direct effects of the variables and doesn't enable to identify combined effect of variables or anything

as such. It should mainly be considered as an informational tool, not an absolute reliable measure of the variables' contributions to the models.


5.1.3 PA data generated

Biomod.PA.data contains the amount of data available after the inner run of the pseudo-absence function. Biomod.PA.sample contains the rows to take from DataBIOMOD to get the data that has been used for the calibration of each species for each PA run.

For example, let's see what data has been used for the calibration of the run PA1 :

```
R code
our.lines <- Biomod.PA.sample$Sp281$PA1
par(mfrow=c(1,2))
level.plot(DataBIOMOD[, "Sp281"], LatLong, title='original data',
           show.scale=F, cex=0.6)
level.plot(DataBIOMOD[our.lines, "Sp281"], LatLong[our.lines,],
           title='PA1', show.scale=F, cex=0.6)
```

original data



PA1



5.2 Objects stored on the hard drive : The Models

Each algorithm (excepted SRE) generates an object storing the different parameterisation, the importance of each variable for the model and other statistics. This output is essential as it allows generating predictions.

These objects, the models themselves, are now stored out of the R workspace directly on the computers' hard disk. They are named after the algorithm used and the species' names, i.e. Sp164_FDA for example. There is also extensions of the names concerning the repetitions and the pseudo-absences runs, so that one of our models will be Sp164_FDA_PA1_rep2.

Back loading the models and having them directly usable is very straightforward : simply use the `load()` function to have the model restored in the R workspace, with the same name plus the directory root. This is also the case with the other outputs stored outside of R (predictions and projections). The syntax is not always handy but easy to pick up :

```
#Example of the GLM
load("models/Sp290_GLM_PA1")
ls()
```

	R code
[1] "BestModelByRoc"	"BestModelByTSS"
[3] "biomodDependencies"	"Biomod.material"
[5] "Biomod.PA.data"	"Biomod.PA.sample"
[7] "DataBIOMOD"	"DataEvalBIOMOD"
[9] "data.used"	"Evaluation.results.Kappa"
[11] "Evaluation.results.Roc"	"Evaluation.results.TSS"
[13] "Expl.Var"	"Expl.Var2"
[15] "Expl.Var3"	"Future1"
[17] "GBM.list"	"GBM.perf"
[19] "i"	"isnulllyweights"
[21] "LatLong"	"missingPackages"
[23] "model"	"myPackages"
[25] "obj"	"our.lines"
[27] "Pred"	"Pred2"
[29] "Pred3"	"PredBestModelByKappa"
[31] "Pred_Sp281"	"Pred_Sp290"
[33] "Pred_Sp290_BinKappa"	"Pred_Sp290_FiltKappa"
[35] "Pred_Sp290_indpdt"	"rand"
[37] "Resp.Var"	"Sp290_GLM_PA1"
[39] "Sp290_RF_PA1"	"Sp.Env"
[41] "store"	"VarImportance"

```
Sp290_GLM_PA1
```

R code

```
Call: glm(formula = Sp290 ~ poly(Var6, 3) + poly(Var7, 3) + poly(Var2,
3) + I(Var1^3) + poly(Var5, 3) + poly(Var3, 2), family = binomial,
data = DataBIOMOD[calib.lines, ], weights = RunWeights[calib.lines])
```

Coefficients:

(Intercept)	poly(Var6, 3)1	poly(Var6, 3)2
-27.5	-138.7	-62.9
poly(Var6, 3)3	poly(Var7, 3)1	poly(Var7, 3)2
-64.1	141.8	-269.4
poly(Var7, 3)3	poly(Var2, 3)1	poly(Var2, 3)2
106.6	615.0	-70.7
poly(Var2, 3)3	I(Var1^3)	poly(Var5, 3)1
-108.7	43.8	-96.8
poly(Var5, 3)2	poly(Var5, 3)3	poly(Var3, 2)1
66.7	32.1	114.8
poly(Var3, 2)2		
-54.1		

Degrees of Freedom: 1772 Total (i.e. Null); 1757 Residual

Null Deviance: 3740

Residual Deviance: 254 AIC: 280

R code

```
summary(Sp290_GLM_PA1)
```

R code

```
Call:
glm(formula = Sp290 ~ poly(Var6, 3) + poly(Var7, 3) + poly(Var2,
3) + I(Var1^3) + poly(Var5, 3) + poly(Var3, 2), family = binomial,
data = DataBIOMOD[calib.lines, ], weights = RunWeights[calib.lines])
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-3.673	0.000	0.000	0.006	3.572

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-27.52	4.03	-6.83	8.8e-12 ***
poly(Var6, 3)1	-138.68	396.79	-0.35	0.72671
poly(Var6, 3)2	-62.87	165.30	-0.38	0.70369
poly(Var6, 3)3	-64.07	39.28	-1.63	0.10289
poly(Var7, 3)1	141.83	134.39	1.06	0.29128
poly(Var7, 3)2	-269.40	47.81	-5.63	1.8e-08 ***
poly(Var7, 3)3	106.58	40.03	2.66	0.00776 **
poly(Var2, 3)1	615.01	315.46	1.95	0.05123 .
poly(Var2, 3)2	-70.70	91.46	-0.77	0.43951
poly(Var2, 3)3	-108.67	32.09	-3.39	0.00071 ***

```

I(Var1^3)          43.80          6.19          7.07  1.5e-12 ***
poly(Var5, 3)1     -96.78          28.28         -3.42  0.00062 ***
poly(Var5, 3)2       66.66          25.21          2.64  0.00819 **
poly(Var5, 3)3       32.14          14.23          2.26  0.02387 *
poly(Var3, 2)1      114.81          28.53          4.02  5.7e-05 ***
poly(Var3, 2)2     -54.14          20.52         -2.64  0.00834 **
---
Signif. codes:  0

```

A series of commands enables you to navigate in the object and to extract usefull information from it. Here are a few example that can be used for all algorithms.

```

#simply type its name
Sp290_GLM_PA1

```

```

R code
Call:  glm(formula = Sp290 ~ poly(Var6, 3) + poly(Var7, 3) + poly(Var2,
3) + I(Var1^3) + poly(Var5, 3) + poly(Var3, 2), family = binomial,
data = DataBIOMOD[calib.lines, ], weights = RunWeights[calib.lines])

```

```

Coefficients:
(Intercept) poly(Var6, 3)1 poly(Var6, 3)2
-27.5        -138.7        -62.9
poly(Var6, 3)3 poly(Var7, 3)1 poly(Var7, 3)2
-64.1         141.8        -269.4
poly(Var7, 3)3 poly(Var2, 3)1 poly(Var2, 3)2
106.6         615.0        -70.7
poly(Var2, 3)3      I(Var1^3) poly(Var5, 3)1
-108.7           43.8        -96.8
poly(Var5, 3)2 poly(Var5, 3)3 poly(Var3, 2)1
66.7            32.1         114.8
poly(Var3, 2)2
-54.1

```

```

Degrees of Freedom: 1772 Total (i.e. Null); 1757 Residual
Null Deviance:          3740
Residual Deviance: 254      AIC: 280

```

```

R code
names(Sp290_GLM_PA1)

```

```

R code
[1] "coefficients"      "residuals"
[3] "fitted.values"     "effects"

```

```

[5] "R"                "rank"
[7] "qr"               "family"
[9] "linear.predictors" "deviance"
[11] "aic"              "null.deviance"
[13] "iter"             "weights"
[15] "prior.weights"    "df.residual"
[17] "df.null"          "y"
[19] "converged"        "boundary"
[21] "model"            "call"
[23] "formula"          "terms"
[25] "data"             "offset"
[27] "control"          "method"
[29] "contrasts"        "xlevels"
[31] "anova"

```

R code

```
str(Sp290_GLM_PA1)
```

R code

List of 31

```

$ coefficients      : Named num [1:16] -27.5 -138.7 -62.9 -64.1 141.8 ...
.. attr(*, "names")= chr [1:16] "(Intercept)" "poly(Var6, 3)1" "poly(Var6, 3)2" "poly(Var6, 3)3" ...
$ residuals         : Named num [1:1773] 588.77 1.11 1 1 -1 ...
.. attr(*, "names")= chr [1:1773] "1" "2" "16" "17" ...
$ fitted.values      : Named num [1:1773] 1.70e-03 9.01e-01 1.00 1.00 2.52e-07 ...
.. attr(*, "names")= chr [1:1773] "1" "2" "16" "17" ...
$ effects            : Named num [1:1773] 4.661 -0.229 3.839 -0.824 -1.76 ...
.. attr(*, "names")= chr [1:1773] "(Intercept)" "poly(Var6, 3)1" "poly(Var6, 3)2" "poly(Var6, 3)3" ...
$ R                  : num [1:16, 1:16] -6.02 0 0 0 0 ...
.. attr(*, "dimnames")=List of 2
.. ..$ : chr [1:16] "(Intercept)" "poly(Var6, 3)1" "poly(Var6, 3)2" "poly(Var6, 3)3" ...
.. ..$ : chr [1:16] "(Intercept)" "poly(Var6, 3)1" "poly(Var6, 3)2" "poly(Var6, 3)3" ...
$ rank               : int 16
$ qr                 :List of 5
..$ qr              : num [1:1773, 1:16] -6.019213 0.04959 0.001657 0.000331 0.000149 ...
.. ..- attr(*, "dimnames")=List of 2
.. .. ..$ : chr [1:1773] "1" "2" "16" "17" ...
.. .. ..$ : chr [1:16] "(Intercept)" "poly(Var6, 3)1" "poly(Var6, 3)2" "poly(Var6, 3)3" ...
..$ rank            : int 16
..$ qraux: num [1:16] 1.01 1.1 1 1 1 ...
..$ pivot: int [1:16] 1 2 3 4 5 6 7 8 9 10 ...
..$ tol            : num 1e-11
.. attr(*, "class")= chr "qr"
$ family             :List of 12
..$ family          : chr "binomial"
..$ link             : chr "logit"
..$ linkfun          :function (mu)

```



```

..$ linkinv :function (eta)
..$ variance :function (mu)
..$ dev.resids:function (y, mu, wt)
..$ aic :function (y, n, mu, wt, dev)
..$ mu.eta :function (eta)
..$ initialize: expression({ if (NCOL(y) == 1) { if (is.factor(y))
..$ validmu :function (mu)
..$ valideta :function (eta)
..$ simulate :function (object, nsim)
..- attr(*, "class")= chr "family"
$ linear.predictors: Named num [1:1773] -6.38 2.21 9.22 12.44 -15.19 ...
..- attr(*, "names")= chr [1:1773] "1" "2" "16" "17" ...
$ deviance : num 254
$ aic : num 280
$ null.deviance : num 3743
$ iter : int 11
$ weights : Named num [1:1773] 1.70e-03 8.91e-02 9.95e-05 3.96e-06 8.05e-07 ...
..- attr(*, "names")= chr [1:1773] "1" "2" "16" "17" ...
$ prior.weights : Named num [1:1773] 1 1 1 1 3.19 ...
..- attr(*, "names")= chr [1:1773] "1" "2" "16" "17" ...
$ df.residual : int 1757
$ df.null : int 1772
$ y : Named num [1:1773] 1 1 1 1 0 0 0 0 0 0 ...
..- attr(*, "names")= chr [1:1773] "1" "2" "16" "17" ...
$ converged : logi TRUE
$ boundary : logi FALSE
$ model : 'data.frame': 1773 obs. of 8 variables:
..$ Sp290 : int [1:1773] 1 1 1 1 0 0 0 0 0 0 ...
..$ poly(Var6, 3): poly [1:1773, 1:3] 0.0465 0.0448 0.0417 0.0344 0.0471 ...
.. ..- attr(*, "dimnames")=List of 2
.. .. ..$ : NULL
.. .. ..$ : chr [1:3] "1" "2" "3"
.. ..- attr(*, "degree")= int [1:3] 1 2 3
.. ..- attr(*, "coefs")=List of 2
.. .. ..$ alpha: num [1:3] 7.9 5.95 8.03
.. .. ..$ norm2: num [1:5] 1 1773 36104 1347227 37019654
.. ..- attr(*, "class")= chr [1:2] "poly" "matrix"
..$ poly(Var7, 3): poly [1:1773, 1:3] 0.0481 0.0468 0.0436 0.0368 0.0464 ...
.. ..- attr(*, "dimnames")=List of 2
.. .. ..$ : NULL
.. .. ..$ : chr [1:3] "1" "2" "3"
.. ..- attr(*, "degree")= int [1:3] 1 2 3
.. ..- attr(*, "coefs")=List of 2
.. .. ..$ alpha: num [1:3] -2.5 -5.66 -3.89
.. .. ..$ norm2: num [1:5] 1.00 1.77e+03 7.73e+04 5.41e+06 2.96e+08
.. ..- attr(*, "class")= chr [1:2] "poly" "matrix"
..$ poly(Var2, 3): poly [1:1773, 1:3] 0.0629 0.0598 0.0544 0.0414 0.064 ...
.. ..- attr(*, "dimnames")=List of 2
.. .. ..$ : NULL

```

```

.. .. .$ : chr [1:3] "1" "2" "3"
.. ..- attr(*, "degree")= int [1:3] 1 2 3
.. ..- attr(*, "coefs")=List of 2
.. .. .$ alpha: num [1:3] 1847 2513 2468
.. .. .$ norm2: num [1:5] 1.00 1.77e+03 1.52e+09 2.56e+15 3.26e+21
.. ..- attr(*, "class")= chr [1:2] "poly" "matrix"
..$ I(Var1^3) :Class 'AsIs' num [1:1773] 0.298 0.438 0.533 0.552 0.123 ...
..$ poly(Var5, 3): poly [1:1773, 1:3] 0.02708 0.03962 0.04306 0.04945 0.00827 ...
.. ..- attr(*, "dimnames")=List of 2
.. .. .$ : NULL
.. .. .$ : chr [1:3] "1" "2" "3"
.. ..- attr(*, "degree")= int [1:3] 1 2 3
.. ..- attr(*, "coefs")=List of 2
.. .. .$ alpha: num [1:3] 179 362 450
.. .. .$ norm2: num [1:5] 1.00 1.77e+03 1.83e+07 5.61e+11 2.02e+16
.. ..- attr(*, "class")= chr [1:2] "poly" "matrix"
..$ poly(Var3, 2): poly [1:1773, 1:2] -0.00208 0.00958 0.01441 0.02074 -0.01698 ...
.. ..- attr(*, "dimnames")=List of 2
.. .. .$ : NULL
.. .. .$ : chr [1:2] "1" "2"
.. ..- attr(*, "degree")= int [1:2] 1 2
.. ..- attr(*, "coefs")=List of 2
.. .. .$ alpha: num [1:2] 798 1470
.. .. .$ norm2: num [1:4] 1.00 1.77e+03 1.84e+08 7.29e+13
.. ..- attr(*, "class")= chr [1:2] "poly" "matrix"
..$ (weights) : num [1:1773] 1 1 1 1 3.19 ...
..- attr(*, "terms")=Classes 'terms', 'formula' length 3 Sp290 ~ poly(Var6, 3) + poly(Var7, 3) + poly(Var2, 3)
.. ..- attr(*, "variables")= language list(Sp290, poly(Var6, 3), poly(Var7, 3), poly(Var2, 3))
.. ..- attr(*, "factors")= int [1:7, 1:6] 0 1 0 0 0 0 0 0 0 1 ...
.. ..- attr(*, "dimnames")=List of 2
.. .. .. .$ : chr [1:7] "Sp290" "poly(Var6, 3)" "poly(Var7, 3)" "poly(Var2, 3)" ...
.. .. .. .$ : chr [1:6] "poly(Var6, 3)" "poly(Var7, 3)" "poly(Var2, 3)" "I(Var1^3)" ...
.. ..- attr(*, "term.labels")= chr [1:6] "poly(Var6, 3)" "poly(Var7, 3)" "poly(Var2, 3)" "I(Var1^3)" ...
.. ..- attr(*, "order")= int [1:6] 1 1 1 1 1 1
.. ..- attr(*, "intercept")= int 1
.. ..- attr(*, "response")= int 1
.. ..- attr(*, ".Environment")=<environment: 0x451cf18>
.. ..- attr(*, "predvars")= language list(Sp290, poly(Var6, 3), coefs = structure(list(alpha = 1847, norm2 = 3.26e+21), type = "double"), poly(Var7, 3), poly(Var2, 3))
.. ..- attr(*, "dataClasses")= Named chr [1:8] "numeric" "nmatrix.3" "nmatrix.3" "nmatrix.3" "numeric" "numeric" "numeric" "numeric"
.. ..- attr(*, "names")= chr [1:8] "Sp290" "poly(Var6, 3)" "poly(Var7, 3)" "poly(Var2, 3)" "I(Var1^3)" "poly(Var6, 3)" "poly(Var7, 3)" "poly(Var2, 3)"
$ call : language glm(formula = Sp290 ~ poly(Var6, 3) + poly(Var7, 3) + poly(Var2, 3), data = dataClasses, weights = weights, method = "OLS")
$ formula :Class 'formula' length 3 Sp290 ~ poly(Var6, 3) + poly(Var7, 3) + poly(Var2, 3)
.. ..- attr(*, ".Environment")=<environment: 0x451cf18>
$ terms :Classes 'terms', 'formula' length 3 Sp290 ~ poly(Var6, 3) + poly(Var7, 3) + poly(Var2, 3)
.. ..- attr(*, "variables")= language list(Sp290, poly(Var6, 3), poly(Var7, 3), poly(Var2, 3))
.. ..- attr(*, "factors")= int [1:7, 1:6] 0 1 0 0 0 0 0 0 0 1 ...
.. ..- attr(*, "dimnames")=List of 2
.. .. .. .$ : chr [1:7] "Sp290" "poly(Var6, 3)" "poly(Var7, 3)" "poly(Var2, 3)" ...
.. .. .. .$ : chr [1:6] "poly(Var6, 3)" "poly(Var7, 3)" "poly(Var2, 3)" "I(Var1^3)" ...

```

```

.. ..- attr(*, "term.labels")= chr [1:6] "poly(Var6, 3)" "poly(Var7, 3)" "poly(Var2, 3)"
.. ..- attr(*, "order")= int [1:6] 1 1 1 1 1 1
.. ..- attr(*, "intercept")= int 1
.. ..- attr(*, "response")= int 1
.. ..- attr(*, ".Environment")=<environment: 0x451cf18>
.. ..- attr(*, "predvars")= language list(Sp290, poly(Var6, 3, coefs = structure(list(alph
.. ..- attr(*, "dataClasses")= Named chr [1:8] "numeric" "nmatrix.3" "nmatrix.3" "nmatrix
.. .. ..- attr(*, "names")= chr [1:8] "Sp290" "poly(Var6, 3)" "poly(Var7, 3)" "poly(Var2,
$ data          : 'data.frame':      1773 obs. of  9 variables:
..$ Var1 : num [1:1773] 0.668 0.76 0.811 0.82 0.497 ...
..$ Var2 : num [1:1773] 4296 4174 3964 3458 4340 ...
..$ Var3 : num [1:1773] 770 928 994 1079 568 ...
..$ Var4 : num [1:1773] 39.3 57.3 66.9 71.4 24.3 ...
..$ Var5 : num [1:1773] 295 349 363 391 215 ...
..$ Var6 : num [1:1773] 16.7 16.4 15.8 14.4 16.9 ...
..$ Var7 : num [1:1773] 10.87 10.51 9.62 7.72 10.39 ...
..$ Sp281: int [1:1773] 0 0 0 0 0 0 0 0 0 0 ...
..$ Sp290: int [1:1773] 1 1 1 1 0 0 0 0 0 0 ...
$ offset          : NULL
$ control          :List of 3
..$ epsilon: num 1e-08
..$ maxit : num 25
..$ trace : logi FALSE
$ method          : chr "glm.fit"
$ contrasts         : NULL
$ xlevels          : Named list()
$ anova           :Classes

```

R code

```

#summary
summary(Sp290_GLM_PA1)

```

R code

```

Call:
glm(formula = Sp290 ~ poly(Var6, 3) + poly(Var7, 3) + poly(Var2,
  3) + I(Var1^3) + poly(Var5, 3) + poly(Var3, 2), family = binomial,
  data = DataBIOMOD[calib.lines, ], weights = RunWeights[calib.lines])

```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-3.673	0.000	0.000	0.006	3.572

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-27.52	4.03	-6.83	8.8e-12 ***
poly(Var6, 3)1	-138.68	396.79	-0.35	0.72671
poly(Var6, 3)2	-62.87	165.30	-0.38	0.70369

```

poly(Var6, 3)3    -64.07      39.28    -1.63    0.10289
poly(Var7, 3)1    141.83     134.39     1.06    0.29128
poly(Var7, 3)2   -269.40      47.81    -5.63    1.8e-08 ***
poly(Var7, 3)3    106.58      40.03     2.66    0.00776 **
poly(Var2, 3)1    615.01     315.46     1.95    0.05123 .
poly(Var2, 3)2   -70.70      91.46    -0.77    0.43951
poly(Var2, 3)3   -108.67      32.09    -3.39    0.00071 ***
I(Var1^3)         43.80       6.19     7.07    1.5e-12 ***
poly(Var5, 3)1   -96.78      28.28    -3.42    0.00062 ***
poly(Var5, 3)2    66.66      25.21     2.64    0.00819 **
poly(Var5, 3)3    32.14      14.23     2.26    0.02387 *
poly(Var3, 2)1   114.81      28.53     4.02    5.7e-05 ***
poly(Var3, 2)2   -54.14      20.52    -2.64    0.00834 **
---
Signif. codes:  0

```

It shows the information stored, like the different variables retained in the final model.

The outputs also give the different coefficient values, the degrees of freedom, the residual deviance and the AIC of the final model. Of course, each model's outputs will not give the same information, as it depends on its specificity.

The next call obtains the anova results and the details of the stepwise procedure type. Note that the independent variables are ranked by their AIC importance.

```

R code
Sp290_GLM_PA1$anova

```

```

R code
Stepwise Model Path
Analysis of Deviance Table

```

Initial Model:

```
Sp290 ~ 1
```

Final Model:

```
Sp290 ~ poly(Var6, 3) + poly(Var7, 3) + poly(Var2, 3) + I(Var1^3) +
      poly(Var5, 3) + poly(Var3, 2)
```

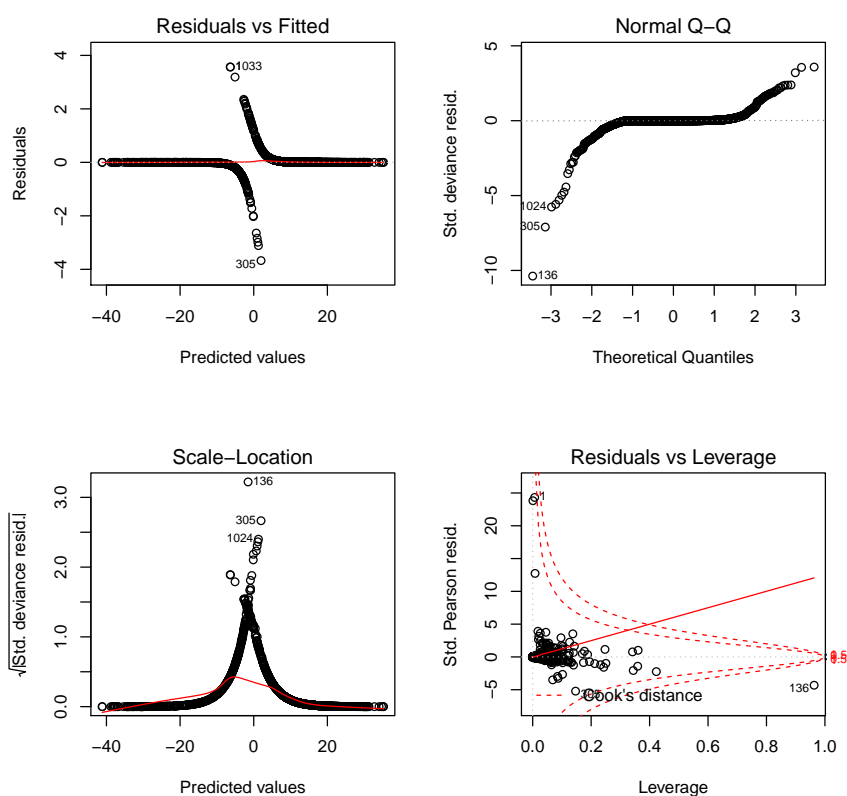
	Step	Df	Deviance	Resid. Df	Resid. Dev	AIC
1				1772	3743.0	3632.7
2	+ poly(Var6, 3)	3	2615.2825	1769	1127.7	1099.9
3	+ Var1	1	494.2727	1768	633.4	624.1
4	+ poly(Var7, 3)	3	285.7301	1765	347.7	355.1

5	+ poly(Var4, 3)	3	23.8282	1762	323.9	337.4
6	+ poly(Var2, 3)	3	22.9058	1759	301.0	321.1
7	+ I(Var1^3)	1	31.0334	1758	269.9	293.4
8	+ poly(Var5, 3)	3	9.5482	1755	260.4	290.2
9	+ poly(Var3, 2)	2	7.6884	1753	252.7	286.8
10	- poly(Var4, 3)	3	0.6872	1756	253.4	281.4
11	- Var1	1	1.0252	1757	254.4	280.3

The function `plot` of R will give the basic and usual outputs for GLM. They are useful but not entirely relevant in the case of the logistic regression.

R code


```
par(mfrow=c(2,2))
plot(Sp290_GLM_PA1)
```



The *gbm* library also provides an experimental diagnostic tool that plots the fitted values versus the actual average values. Uses gam to estimate $E(y|p)$. Well-calibrated predictions imply that $E(y|p) = p$. The plot also includes a pointwise 95 band.

This method can be applied to all models to visualise the relative goodness of fit of the model. The function requires the observed presence-absence of the selected species and the predictions. Hence, you will need to load the predictions for this.

```
library(gbm)
load("pred/Pred_Sp290")
#let's store the data that was used for calibration of the
#first PA run for Sp290 to simplify the code
data.used <- DataBIOMOD[Biomod.PA.sample$Sp290$PA1,"Sp290"]
calibrate.plot(data.used, Pred_Sp290[, "GLM", 1, 1]/1000)
```



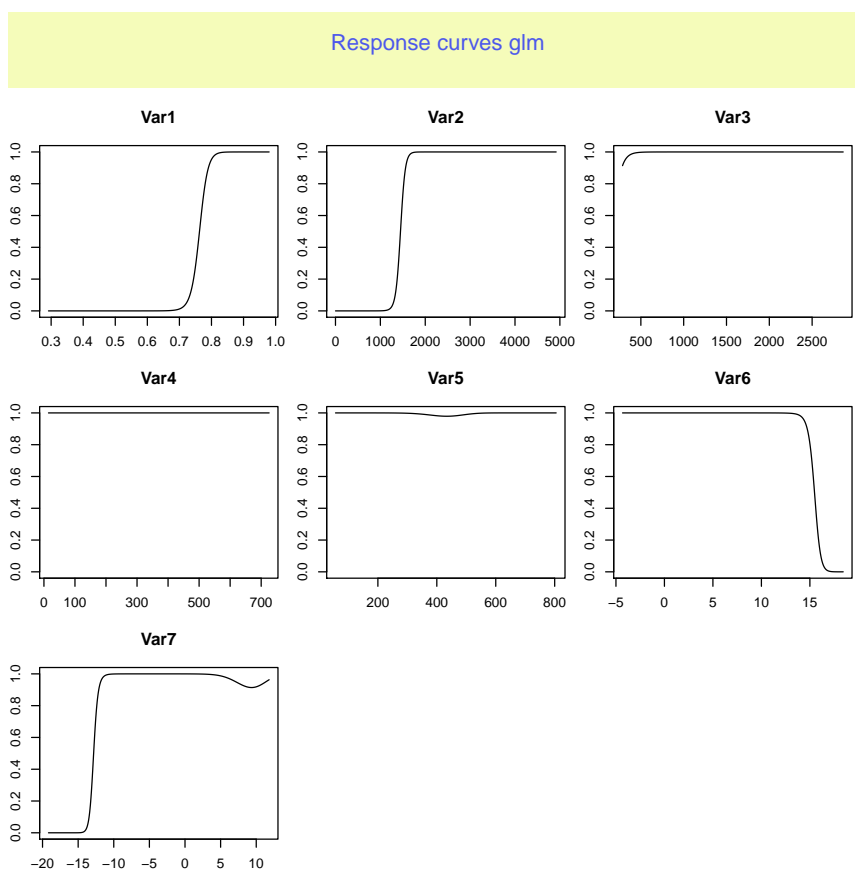
5.2.1 Response Curves

BIOMOD allows plotting the response curves of every model in the good scale. The *response.plot* function must be used to this matter. This function requires a model and a related set of variables to plot the response curves.

Here are two examples of the GLM and RF for the first species modelled. You need to load the model, type its name in the first argument, then give the variables for which you want to see the curves. Note the you can choose to only show some of the variables with the *show.variables* argument.

R code

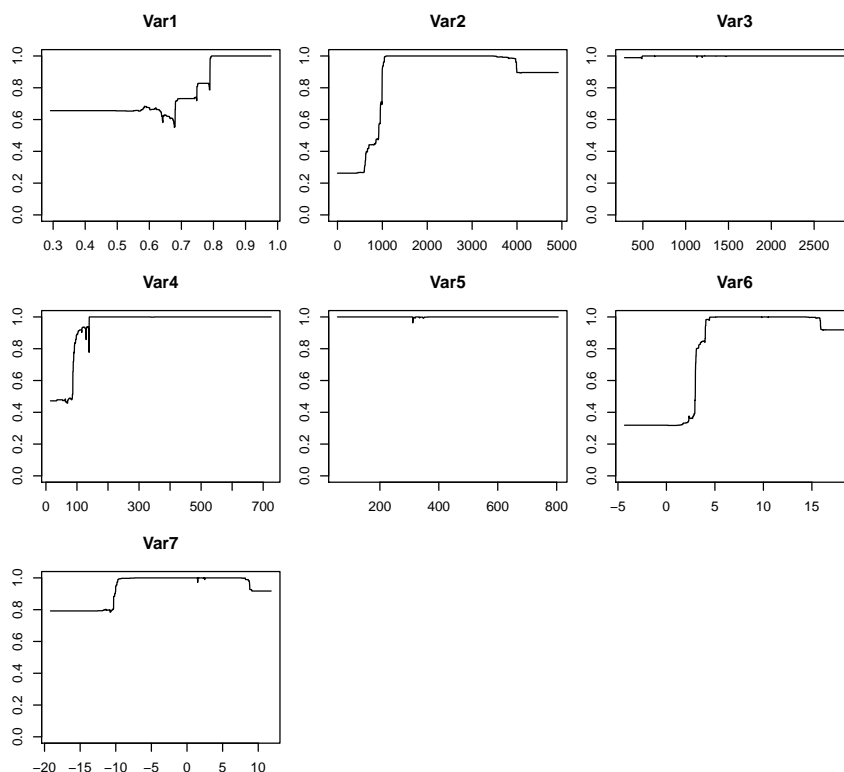
```
#this one has already been loaded in a prior call
response.plot(Sp290_GLM_PA1, Expl.Var)
```



R code

```
load("models/Sp290_RF_PA1")
response.plot(Sp290_RF_PA1, Expl.Var)
```

Response curves randomForest



The response curves are generated following this calculation : N-1 variables are held constant at their mean value whilst the variable of interest contains 100 points varying across the maximum and the minimum of its range. Variation in predictions, made to these 100 cells, only reflects the effects of the one selected variable. Thus, a plot of these predictions allows visualisation of the modelled response to the variable of interest, contingent on the other variables being held constant. This is done subsequently for all the selected variables.

In our examples, the variable Var4 doesn't seem to have a great influence for the GLM (very few variations in the prediction staying close to 1) when it shows a non negligible influence in the predictions of the RF.

These results are interesting when put together with the VarImportance results. They show that Var5 which shows variability from one model to another doesn't has a high importance for most of the models. In contrast, the variable Var6 which is consistent accross GLM and RF has a big influence on the models. This variable is surely connected with the presence/absence of species 290 and the response plots shows this relationship.

5.3 Objects stored on the hard drive : The Predictions

The predictions made by each model on the original dataset are stored inside the *pred* folder. They are stored independently for each species in an object following a 'Pred.Speciesname' logic and contains the probability of occurrence (habitat suitability index) for each run (if several runs) of the selected models. The same objects are produced for the independent data (if any) and the same logic is respected for the projections.

NOTE: for calculation and memory storage purposes, this index is on a scale between 0 and 1000. To obtain a true probability of occurrence, rescaled between 0 and 1, simply divide each value by a thousand.

R code

```
load("pred/Pred_Sp290")
```

The trick is that these objects are no longer matrices but arrays (multiple dimensions) with 4 dimensions. The dimensions can be visualised as follows :

The first two build up a matrix where each column is the prediction of one of the models. The number of rows corresponds to the amount of data used for building those models.

R code

```
dim(Pred_Sp290)
```

R code

```
[1] 1773    9    4    1
```

R code


```
Pred_Sp290[1:20,,1,1]
```

R code

	ANN	CTA	GAM	GBM	GLM	MARS	FDA	RF	SRE
1	61	39	142	412	1	96	35	745	0
2	61	987	777	751	901	974	983	929	0
3	61	987	978	766	999	997	984	988	0
4	986	987	984	766	999	996	984	996	0
5	61	39	1	75	0	3	33	0	0
6	61	39	6	75	0	3	33	0	0
7	61	39	1	75	0	7	33	0	0
8	61	39	9	75	0	3	33	0	0

9	61	39	1	75	0	9	33	0	0
10	61	39	1	75	0	7	33	0	0
11	61	39	19	75	0	4	33	0	0
12	61	39	1	75	0	9	33	0	0
13	61	39	12	76	0	3	33	0	0
14	61	39	90	79	86	4	33	18	0
15	985	987	995	766	999	998	984	996	0
16	986	987	997	767	999	999	984	997	0
17	989	987	999	906	999	999	984	996	0
18	988	987	999	842	999	999	984	994	0
19	989	987	997	911	999	998	984	1000	0
20	989	987	998	887	999	999	984	1000	0

Now, the third dimensions consists of a collection of 2-D matrices, one behind another, corresponding to the prediction produced by each repetition. The minimum for this dimension is 1. Considering that BIOMOD always produces a final model calibrated with 100% of the data given, the length of this third dimension is the value of the NbRunEval argument + 1. For example, with NbRunEval=10, you have 11 layers.



Note that the first layer is always the final model, then come the repetitions.

R code

```
#the final model
Pred_Sp290[1:15,,1,1]
```

	ANN	CTA	GAM	GBM	GLM	MARS	FDA	<i>R code</i>	
								RF	SRE
1	61	39	142	412	1	96	35	745	0
2	61	987	777	751	901	974	983	929	0
3	61	987	978	766	999	997	984	988	0
4	986	987	984	766	999	996	984	996	0
5	61	39	1	75	0	3	33	0	0
6	61	39	6	75	0	3	33	0	0
7	61	39	1	75	0	7	33	0	0

8	61	39	9	75	0	3	33	0	0
9	61	39	1	75	0	9	33	0	0
10	61	39	1	75	0	7	33	0	0
11	61	39	19	75	0	4	33	0	0
12	61	39	1	75	0	9	33	0	0
13	61	39	12	76	0	3	33	0	0
14	61	39	90	79	86	4	33	18	0
15	985	987	995	766	999	998	984	996	0

R code

```
#the first repetition model
Pred_Sp290[1:15,,2,1]
```

	ANN	CTA	GAM	GBM	GLM	MARS	FDA	R code	
								RF	SRE
1	31	319	103	376	6	312	33	708	0
2	31	319	598	616	886	990	985	821	0
3	64	319	950	652	999	998	985	934	0
4	580	319	971	653	999	999	985	926	0
5	31	28	0	74	0	2	31	0	0
6	31	28	3	74	0	2	31	1	0
7	31	28	0	74	0	2	31	2	0
8	31	28	5	74	0	2	31	0	0
9	31	28	0	74	0	2	31	1	0
10	31	28	0	74	0	2	31	1	0
11	31	28	14	74	0	2	31	0	0
12	31	28	0	74	0	2	31	0	0
13	31	28	6	75	0	2	31	0	0
14	31	28	91	82	90	3	31	14	0
15	402	319	989	655	999	999	985	952	0

R code

```
#the second repetition model
Pred_Sp290[1:15,,3,1]
```

	ANN	CTA	GAM	GBM	GLM	MARS	FDA	R code	
								RF	SRE
1	81	36	187	370	0	175	36	410	0
2	925	987	855	757	618	992	983	889	0
3	991	987	988	801	999	998	984	974	0
4	992	987	991	802	999	998	984	984	0
5	45	36	0	74	0	2	32	0	0
6	47	36	1	74	0	4	32	2	0
7	45	36	0	74	0	2	32	0	0
8	47	36	3	74	0	4	32	1	0

```

9  46 36  0 74  0  2 32  0  0
10 46 36  0 74  0  2 32  1  0
11 50 36 12 74  0  6 32  0  0
12 47 36  0 74  0  2 32  0  0
13 46 36  3 74  0  5 32  0  0
14 53 36 96 84 178 13 32 29  0
15 997 987 998 802 999 999 984 993  0


```

```

# R code
#and so on...

```

The fourth dimension represents the number of pseudo-absences repetitions that have been made. In the case where NbRepPA=0, the dimension is simply 1 (not 0).



You will never visualise it this way with R though. It is just an abstract view of how it is sorted. Some usefull functions for not getting lost are `dim()` and `dimnames()`. The first one gives you the number of layers for each dimension, the second will give you their names respectively.

```

# R code
load("pred/Pred_Sp281")
dim(Pred_Sp281)

```

```

# R code
[1] 1392  9  4  2

```

```

# R code
#dimnames(Pred_Sp281)
#you can avoid having the rownames to be printed in the console as they
#are generally not very usefull
dimnames(Pred_Sp281)[-1]

```

```

R code
[[1]]
[1] "ANN" "CTA" "GAM" "GBM" "GLM" "MARS" "FDA" "RF"
[9] "SRE"

[[2]]
[1] "total.data" "rep1" "rep2" "rep3"

[[3]]
[1] "PA1" "PA2"

```

For instance, we examine the probability of occurrence of the first species, modelled with CTA. Here we just display 20 rows (or sites) in the middle.

```

R code
#if you don't inform the 3rd and 4th dimension (you still need commas), you will have all c
#at once in a matrix.
load("pred/Pred_Sp281")
Pred_Sp281[281:300,"CTA",,]

```

```

R code
, , PA1

      total.data rep1 rep2 rep3
281      1000  958 1000  997
282      993  958  994  997
283      993  958  994  997
284      993  958  994  997
285      946  836   12  933
286     1000  958 1000  997
287     1000  836 1000  919
288      792  836  792  919
289      993  958  994  997
290      993  958  994  997
291      993  958  994  997
292      993  958  994  629
293      993  958  994  629
294         0  958  910  997
295         0  958    0  997
296         0  958    0    0
297         0  958    0  629
298         0  958    0  629
299      993  958  994    0
300         0  958    0    0

, , PA2

```


	<i>total.data</i>	<i>rep1</i>	<i>rep2</i>	<i>rep3</i>
281	967	979	990	969
282	998	979	990	969
283	922	1000	968	827
284	967	976	968	946
285	1000	976	968	946
286	1000	976	968	946
287	800	979	990	969
288	998	979	990	969
289	998	979	990	969
290	800	979	990	969
291	998	979	656	86
292	800	772	990	969
293	800	772	990	969
294	203	0	990	969
295	203	0	990	969
296	967	976	656	86
297	967	976	656	86
298	203	979	990	969
299	203	0	656	809
300	800	979	990	969

Note that because there is a random selection of the data for calibration, you will end up with slightly different values on these example runs.

Because we have chosen to run the models with pseudo-absence data, plotting the partial predictions is not very convinient. We will plot instead the values of the fake independent data (which is just the full original dataset) for the GLM.

R code

```
load("pred/Pred_Sp290_indpdt")
level.plot(Pred_Sp290_indpdt[, "GLM", 1, 1], LatLong, title='sp290_GLM_indpdt', cex=0.8)
```



Note that the independent predictions are only made on the final 100% model and not on the repetitions. To check it :

`Pred_Sp290_indpdt[1:10,,]` R code

`, , total.data` R code

	ANN	CTA	GAM	GBM	GLM	MARS	FDA	RF	SRE
1	61	39	142	412	1	96	35	745	0
2	61	987	777	751	901	974	983	929	0
3	61	987	669	698	674	946	973	782	0
4	61	39	2	75	0	2	33	6	0
5	61	39	1	75	0	2	33	6	0
6	61	39	2	75	0	2	33	0	0
7	61	39	1	75	0	2	33	0	0
8	61	39	18	79	0	5	33	12	0
9	61	39	7	75	0	3	33	4	0
10	61	39	25	79	0	3	33	10	0

, , rep1

	ANN	CTA	GAM	GBM	GLM	MARS	FDA	RF	SRE
1	NA	NA	NA	NA	NA	NA	NA	NA	NA
2	NA	NA	NA	NA	NA	NA	NA	NA	NA
3	NA	NA	NA	NA	NA	NA	NA	NA	NA
4	NA	NA	NA	NA	NA	NA	NA	NA	NA
5	NA	NA	NA	NA	NA	NA	NA	NA	NA
6	NA	NA	NA	NA	NA	NA	NA	NA	NA
7	NA	NA	NA	NA	NA	NA	NA	NA	NA
8	NA	NA	NA	NA	NA	NA	NA	NA	NA
9	NA	NA	NA	NA	NA	NA	NA	NA	NA
10	NA	NA	NA	NA	NA	NA	NA	NA	NA

, , rep2

	ANN	CTA	GAM	GBM	GLM	MARS	FDA	RF	SRE
1	NA	NA	NA	NA	NA	NA	NA	NA	NA
2	NA	NA	NA	NA	NA	NA	NA	NA	NA
3	NA	NA	NA	NA	NA	NA	NA	NA	NA
4	NA	NA	NA	NA	NA	NA	NA	NA	NA
5	NA	NA	NA	NA	NA	NA	NA	NA	NA
6	NA	NA	NA	NA	NA	NA	NA	NA	NA
7	NA	NA	NA	NA	NA	NA	NA	NA	NA
8	NA	NA	NA	NA	NA	NA	NA	NA	NA
9	NA	NA	NA	NA	NA	NA	NA	NA	NA
10	NA	NA	NA	NA	NA	NA	NA	NA	NA

, , rep3

	ANN	CTA	GAM	GBM	GLM	MARS	FDA	RF	SRE
1	NA	NA	NA	NA	NA	NA	NA	NA	NA
2	NA	NA	NA	NA	NA	NA	NA	NA	NA
3	NA	NA	NA	NA	NA	NA	NA	NA	NA
4	NA	NA	NA	NA	NA	NA	NA	NA	NA
5	NA	NA	NA	NA	NA	NA	NA	NA	NA
6	NA	NA	NA	NA	NA	NA	NA	NA	NA
7	NA	NA	NA	NA	NA	NA	NA	NA	NA
8	NA	NA	NA	NA	NA	NA	NA	NA	NA
9	NA	NA	NA	NA	NA	NA	NA	NA	NA
10	NA	NA	NA	NA	NA	NA	NA	NA	NA

5.3.1 Transforming the predictions on the original dataset

It might be useful to extract the presence/absence predictions. To do so, use the *CurrentPred()* function by switching *BinRoc*, *BinKappa* and/or *BinTSS* to TRUE and each probability of occurrence will be transformed into pres-

ence and absence using the cutoff maximising the models accuracy according to Roc, Kappa or TSS. You can also selecting the *FiltRoc*, *FiltKappa* and/or *FiltTSS* options. That will result in creating a new table where probabilities lower than corresponding optimised cutoff are set to 0 and those upper keep their value.

R code

```
CurrentPred(GLM=T, GBM=T, GAM=T, CTA=T, ANN=T, SRE=T,
            FDA=T, MARS=F, RF=T, BinRoc=T, BinKappa=T,
            BinTSS=T, FiltKappa=T)
```

New objects are created for each species containing the predictions in binary and or filtered format using the thresholds produced by the evaluation technics : `Pred_Sp290_BinRoc`, `Pred_Sp290_BinKappa`, `Pred_Sp290_BinTSS`, `Pred_Sp290_FiltKappa`, and so on.

R code

```
load("pred/Pred_Sp290")
load("pred/Pred_Sp290_BinKappa")
load("pred/Pred_Sp290_FiltKappa")
Pred_Sp290[260:270,,1,1]
```

R code

	ANN	CTA	GAM	GBM	GLM	MARS	FDA	RF	SRE
260	979	987	998	931	999	998	984	1000	0
261	989	987	987	929	991	997	984	1000	1000
262	979	987	935	932	953	998	984	1000	0
263	989	987	944	930	954	998	984	1000	1000
264	989	987	961	930	954	998	984	1000	1000
265	989	987	999	930	999	997	984	1000	1000
266	989	987	999	930	999	997	984	993	1000
267	989	987	999	930	999	997	984	1000	1000
268	989	987	999	930	999	997	984	1000	1000
269	989	987	999	930	999	997	984	996	1000
270	989	987	998	930	999	997	984	1000	1000

R code

```
Pred_Sp290_BinKappa[260:270,,1,1]
```

R code

	ANN	CTA	GAM	GBM	GLM	MARS	FDA	RF	SRE
260	1	1	1	1	1	NA	1	1	0
261	1	1	1	1	1	NA	1	1	1

262	1	1	1	1	1	NA	1	1	0
263	1	1	1	1	1	NA	1	1	1
264	1	1	1	1	1	NA	1	1	1
265	1	1	1	1	1	NA	1	1	1
266	1	1	1	1	1	NA	1	1	1
267	1	1	1	1	1	NA	1	1	1
268	1	1	1	1	1	NA	1	1	1
269	1	1	1	1	1	NA	1	1	1
270	1	1	1	1	1	NA	1	1	1

R code

```
Pred_Sp290_FiltKappa[260:270,,1,1]
```

	ANN	CTA	GAM	GBM	GLM	MARS	FDA	<i>R code</i>	RF	SRE
260	979	987	998	931	999	NA	984	1000		0
261	989	987	987	929	991	NA	984	1000	1000	
262	979	987	935	932	953	NA	984	1000		0
263	989	987	944	930	954	NA	984	1000	1000	
264	989	987	961	930	954	NA	984	1000	1000	
265	989	987	999	930	999	NA	984	1000	1000	
266	989	987	999	930	999	NA	984	993	1000	
267	989	987	999	930	999	NA	984	1000	1000	
268	989	987	999	930	999	NA	984	1000	1000	
269	989	987	999	930	999	NA	984	996	1000	
270	989	987	998	930	999	NA	984	1000	1000	

R code

5.3.2 Identifying the best model

In our example, we could compare all the models we run for the different species using the three different evaluation methods available. The function *PredictionBestModel* also transforms the probabilities into the presence/absence and filtered formats.

R code

```
PredictionBestModel(GLM=T,GBM=T, GAM=T, CTA=T, ANN=T, FDA=T,
                     MARS=F, RF=T, SRE=T, method='all',
                     Bin.trans = T, Filt.trans = T)
```

Multimodel comparison according to the TSS statistic:

```

load("pred/BestModelByTSS")
BestModelByTSS

```

```

$Sp281
      Best.Model Cross.validation indepdt.data
PA1          RF          0.947          0.876
PA1_rep1      RF          0.944          none
PA1_rep2      ANN          0.952          none
PA1_rep3      RF          0.955          none
PA2          RF          0.952          0.871
PA2_rep1      RF          0.940          none
PA2_rep2      RF          0.945          none
PA2_rep3      RF          0.972          none
      total.score Cutoff Sensitivity Specificity
PA1          1.0000  340.0      100.00      100.0
PA1_rep1      0.9889  410.0      99.49       99.4
PA1_rep2      0.9394  431.6      96.94       97.0
PA1_rep3      0.9899  420.0      99.49       99.5
PA2          1.0000  390.0      100.00      100.0
PA2_rep1      0.9868  450.0      98.98       99.7
PA2_rep2      0.9843  380.0      99.23       99.2
PA2_rep3      0.9939  490.0      99.49       99.9

```

```

$Sp290
      Best.Model Cross.validation indepdt.data
PA1          RF          0.978          0.784
PA1_rep1      GAM          0.981          none
PA1_rep2      RF          0.978          none
PA1_rep3      RF          0.981          none
      total.score Cutoff Sensitivity Specificity
PA1          1.0000  350.0      100.00      100.00
PA1_rep1      0.9666  409.6      98.07       98.58
PA1_rep2      0.9933  710.0      99.33      100.00
PA1_rep3      0.9962  390.0      99.85       99.76

```

The RF comes out first almost each time, let's switch it off : Multimodel comparison according to the TSS statistic:

```

PredictionBestModel(GLM=T,GBM=T, GAM=T, CTA=T, ANN=T, FDA=T,
                    MARS=F, RF=F, SRE=T, method='all',
                    Bin.trans = T, Filt.trans = T)
load("pred/BestModelByTSS")
BestModelByTSS

```

R code

```

$Sp281
      Best.Model Cross.validation indepth.data
PA1      GBM      0.922      0.776
PA1_rep1 GBM      0.919      none
PA1_rep2 ANN      0.952      none
PA1_rep3 GLM      0.934      none
PA2      GBM      0.918      0.778
PA2_rep1 GBM      0.889      none
PA2_rep2 ANN      0.912      none
PA2_rep3 FDA      0.962      none
      total.score Cutoff Sensitivity Specificity
PA1      0.9388 474.04      98.98      94.9
PA1_rep1 0.9316 568.86      97.96      95.2
PA1_rep2 0.9394 431.64      96.94      97.0
PA1_rep3 0.9446 499.50      97.96      96.5
PA2      0.9436 577.40      97.96      96.4
PA2_rep1 0.9279 527.54      97.19      95.6
PA2_rep2 0.9306 495.00      97.96      95.1
PA2_rep3 0.9222 70.32      95.92      96.3

$Sp290
      Best.Model Cross.validation indepth.data
PA1      ANN      0.971      0.658
PA1_rep1 ANN      0.981      none
PA1_rep2 GBM      0.970      none
PA1_rep3 GLM      0.974      none
      total.score Cutoff Sensitivity Specificity
PA1      0.9385 358.0      97.63      96.22
PA1_rep1 0.9666 409.6      98.07      98.58
PA1_rep2 0.9732 457.4      98.74      98.58
PA1_rep3 0.9645 459.5      97.63      98.82

```

Multimodel comparison according to the ROC:

R code

```

load("pred/BestModelByRoc")
BestModelByRoc

```

R code

```

$Sp281
      Best.Model Cross.validation indepth.data
PA1      GBM      0.99      0.941
PA1_rep1 GBM      0.988      none
PA1_rep2 GBM      0.991      none
PA1_rep3 GBM      0.992      none
PA2      GBM      0.992      0.941
PA2_rep1 GBM      0.988      none
PA2_rep2 GBM      0.991      none

```

	FDA	0.998	none
	total.score	Cutoff	Sensitivity Specificity
PA2_rep3	0.996	603.784	96.173 96.3
PA1	0.994	614.976	96.173 96.2
PA1_rep1	0.993	626.373	96.684 96.7
PA1_rep2	0.995	589.889	96.173 96.2
PA1_rep3	0.996	616.454	96.429 96.4
PA2	0.995	583.706	96.173 96.1
PA2_rep1	0.995	587.232	95.918 96
PA2_rep2	0.987	66.488	96.173 96.1
PA2_rep3			

\$Sp290

	Best.Model	Cross.validation	indept.data
PA1	GBM	0.998	0.914
PA1_rep1	GBM	0.998	none
PA1_rep2	GBM	0.997	none
PA1_rep3	GBM	0.998	none
	total.score	Cutoff	Sensitivity Specificity
PA1	0.999	450.142	98.593 98.582
PA1_rep1	0.999	460.802	98.593 98.582
PA1_rep2	0.999	476.336	98.593 98.582
PA1_rep3	0.998	408.591	98 98.109

Multimodel predictions according to the Kappa statistic

```

R code
load("pred/PredBestModelByKappa")
PredBestModelByKappa[740:750,,1]

```

	PA1	PA1_rep1	PA1_rep2	PA1_rep3	PA2	PA2_rep1	PA2_rep2
740	71	71	0	0	34	23	72
741	71	71	0	0	34	23	72
742	71	71	0	0	34	23	72
743	71	71	0	0	34	23	72
744	71	71	0	0	34	23	72
745	71	71	1	2	34	23	72
746	71	71	1	2	34	23	71
747	71	71	1	0	34	23	75
748	71	71	0	0	34	23	72
749	71	72	1	0	35	23	94
750	71	72	0	0	34	23	72
PA2_rep3							
740		32					
741		32					
742		32					

743	32
744	32
745	33
746	32
747	33
748	32
749	32
750	32

6 Uncertainty analysis

6.1 Models' projection

For all the models currently implemented, BIOMOD is able to project potential distributions of species or land-use classes for other areas, other resolutions or other times. BIOMOD does not utilise the geographical coordinates nor does it perform a re-ordering of the data for making projections. **The user must ensure** that all datasets are kept in the same order in order to allow unmistakable comparisons between observed and predicted maps.

To make the projections, use the function *Projection*.

The syntax is very similar to previous functions. First add the new data (e.g. climate change scenario), then the prefix name of the output (Proj.name), and then the models for which the projections have to be made.

The Proj.name argument is very important as it will be used to store the results and also used by other functions to reload this data. The *Projection* function will create a directory using that name. In our case, it will produce "proj.Future1" next to "pred" and "models" in the working directory. A directory is created for each run of the function with a different scenario.

```
#load the example dataset : future scenario 1
data(Future1)
head(Future1)
Projection(Proj = Future1[,4:10], Proj.name='Future1',
          GLM = T, GBM = T, GAM = T, CTA = T, ANN = T,
          SRE = T, quant=0.025, MARS = T, RF = T,
          BinRoc = T, BinKappa = T, BinTSS = T, FiltRoc = T,
          FiltKappa = T, FiltTSS = T, repetition.models=T)
save.image('RUN.RData')
```

Let's check the future projections made for this scenario :

```
load('RUN.RData')
load("proj.Future1/Proj_Future1_Sp290")
dim(Proj_Future1_Sp290)
```

```
[1] 2264    9    4    1
```

```

dimnames(Proj_Future1_Sp290)[-1]

```

```

[[1]]
[1] "ANN"  "CTA"  "GAM"  "GBM"  "GLM"  "MARS" "FDA"  "RF"
[9] "SRE"

[[2]]
[1] "total.data" "rep1"      "rep2"      "rep3"

[[3]]
[1] "PA1"

```

```

Proj_Future1_Sp290[740:750,,1,1]

```

```

      ANN CTA  GAM GBM  GLM  MARS  FDA  RF  SRE
740 989 987 999 930 999  998 984 1000 1000
741 986 987 999 930 999  997 984  992    0
742 989 987 999 930 999  997 984 1000 1000
743 985 987 999 930 999  996 984  985    0
744 979 987 997 929 999  999 984  997 1000
745 981 987 999 930 999  998 984 1000 1000
746 979 916 710 828 853  956 983  956    0
747 989 987 982 930 999  991 984 1000 1000
748 989 987 999 930 999  997 984 1000 1000
749 989 987 999 930 999  998 984 1000 1000
750 989 987 999 930 999  997 984 1000 1000

```

```

load("proj.Future1/Proj_Future1_Sp290_BinRoc")
Proj_Future1_Sp290_BinRoc[740:750,,1,1]

```

```

      ANN CTA  GAM GBM  GLM  MARS  FDA  RF  SRE
740    1    1    1    1    1    1    1    1    1
741    1    1    1    1    1    1    1    1    0
742    1    1    1    1    1    1    1    1    1
743    1    1    1    1    1    1    1    1    0
744    1    1    1    1    1    1    1    1    1
745    1    1    1    1    1    1    1    1    1
746    1    1    1    1    1    1    1    1    0

```

747	1	1	1	1	1	1	1	1	1
748	1	1	1	1	1	1	1	1	1
749	1	1	1	1	1	1	1	1	1
750	1	1	1	1	1	1	1	1	1

We also have them in binary and filtered format which have been directly produced by the Projection function.

Compare the projections produced with original data

```

R code
multiple.plot(cbind(Sp.Env[, 'Sp290'], Proj_Future1_Sp290[,1:8,1,1]), LatLong, cex=0.73)

```

```

R code
load("proj.Future1/Proj_Future1_Sp281")
#PA1 et PA2
multiple.plot(cbind(PA1=Sp.Env[, 'Sp281'], PA2=Sp.Env[, 'Sp281'], Proj_Future1_Sp281[,1:8,1,1]), LatLong, cex=0.73)

```

```

R code
#repetitions
multiple.plot(cbind(full=Sp.Env[, 'Sp281'], Proj_Future1_Sp281[,1:8,1,2]), LatLong, cex=0.73)
multiple.plot(cbind(rep1=Sp.Env[, 'Sp281'], Proj_Future1_Sp281[,1:8,2,2]), LatLong, cex=0.73)
multiple.plot(cbind(rep2=Sp.Env[, 'Sp281'], Proj_Future1_Sp281[,1:8,3,2]), LatLong, cex=0.73)
multiple.plot(cbind(rep3=Sp.Env[, 'Sp281'], Proj_Future1_Sp281[,1:8,4,2]), LatLong, cex=0.73)

```

So we have here 9x4 projections for each PA run, which gives 72 projections per future scenario (2 PA runs). So in total : 144 projections.

6.2 Ensemble Forecasting

Several approaches are available for combining ensembles of models in BIOMOD. Here is an example of the use of the *Ensemble.Forecasting* function as well as some details of the different strategies:

Four straightforward means of 'committee averaging' (giving the same weight to all the elements) are done across all the models for each run:

- on the probabilities
- on the binary projection according to the Roc method
- on the binary projection according to the Kappa method
- on the binary projection according to the TSS method

A weighted approach is also available that ranks the models using their evaluation score.

Making a mean on the 0-1 projections gives some sort of probability of occurrence. For example, for a given site and with the TSS method, 6 projections give a "1" and 2 give a "0". The mean will be 0.75. It is extracted from binary projection and it is therefore not possible to determine a prior threshold. Conversion into binary is nevertheless possible (see *binary* below).

The median value is also calculated on the probabilities given by the models. It is considered to be more reliable because it is less influenced by extreme values.

Some options:

repetition.models: You can choose to switch on or off the repetition models. If selected, the function will calculate the ensemble forecasts for each run and generate a final one which produces a general ensemble forecast across all the runs for each method. This total consensus is done inconsistently of this argument being set to TRUE or FALSE.

weight.method: the method for ranking the models according to their predictive performance. The *decay* gives the relative importance of the weights. The default weight decay is 1.6; See the example below.

models	GAM	GBM	GLM	ANN	RF	MARS	CTA	FDA
score with Roc	0.96	0.92	0.90	0.88	0.87	0.75	0.72	0.68
decay of 1	0.125	0.125	0.125	0.125	0.125	0.125	0.125	0.125

decay of 1.2	0.217	0.181	0.151	0.126	0.105	0.087	0.073	0.061
decay of 1.6	0.384	0.240	0.150	0.094	0.059	0.037	0.023	0.014
decay of 2	0.502	0.251	0.125	0.063	0.031	0.016	0.008	0.004

You can type in any value (it has however to be higher than 1) depending on the strength of discrimination that you want. A decay of 1 is equivalent to a committee averaging (i.e. same weights given to all elements).

final.model.out: set to True if you want the total ensemble to be build with the final models taken into account.

qual.th: enables to switch off the models under a certain evaluation score. This will be applied to all models on all species. This option is usefull if you think some of your models are really bad for your study case.

compress: logical or character string specifying whether saving to a named file is to use compression. FALSE corresponds to no compression, and character strings "gzip", or "xz" specify the type of compression. See ?save for more details. Default is "xz". Note that compression may be a long task so you can switch it off if you are more interesting in saving time than in saving space.

```
R code
Ensemble.Forecasting(Proj.name= "Future1", weight.method='Roc',
                     PCA.median=T, binary=T, bin.method='Roc',
                     Test=F, decay=1.6, repetition.models=T,
                     final.model.out=FALSE, qual.th=0, compress="xz")
```

R code

Sp281

Sp290

consensus_Future1_results

\$Sp281

\$Sp281\$weights

	ANN	CTA	GAM	GBM	GLM	MARS	FDA
PA1	0.0297	0.0143	0.0297	0.1500	0.0762	0.3120	0.0762
PA1_rep1	0.0229	0.0143	0.0586	0.1500	0.0366	0.3839	0.0937
PA1_rep2	0.0297	0.0143	0.0297	0.1219	0.1219	0.2400	0.0586
PA1_rep3	0.0366	0.0143	0.0229	0.1500	0.0937	0.2400	0.0586
PA2	0.0366	0.0229	0.0586	0.2400	0.0143	0.1500	0.0937
PA2_rep1	0.0366	0.0229	0.0586	0.2400	0.0143	0.1500	0.0937
PA2_rep2	0.0366	0.0143	0.0586	0.2400	0.0229	0.1500	0.0937
PA2_rep3	0.0476	0.0229	0.0476	0.0937	0.0143	0.3120	0.1500

RF SRE

```

PA1      0.3120  0
PA1_rep1 0.2400  0
PA1_rep2 0.3839  0
PA1_rep3 0.3839  0
PA2      0.3839  0
PA2_rep1 0.3839  0
PA2_rep2 0.3839  0
PA2_rep3 0.3120  0

```

```
$Sp281$PCA.median
```

```
model.selected
```

```

PA1      "MARS"
PA1_rep1 "GBM"
PA1_rep2 "GLM"
PA1_rep3 "GLM"
PA2      "MARS"
PA2_rep1 "RF"
PA2_rep2 "GBM"
PA2_rep3 "FDA"

```

```
$Sp281$thresholds
```

	PA1	PA1_rep1	PA1_rep2	PA1_rep3	PA2
prob.mean	496.8	511.7	451.3	451.6	502.4
prob.mean.weighted	465.1	387.9	370.9	402.7	519.0
median	572.6	609.4	498.2	499.7	576.0
Roc.mean	500.0	500.0	500.0	500.0	500.0
Kappa.mean	500.0	500.0	500.0	500.0	500.0
TSS.mean	500.0	500.0	500.0	500.0	500.0

	PA2_rep1	PA2_rep2	PA2_rep3
prob.mean	470.4	491.7	454.0
prob.mean.weighted	406.8	421.2	332.1
median	543.2	584.3	487.9
Roc.mean	500.0	500.0	500.0
Kappa.mean	500.0	500.0	500.0
TSS.mean	500.0	500.0	500.0

```
$Sp290
```

```
$Sp290$weights
```

	ANN	CTA	GAM	GBM	GLM	MARS	FDA
PA1	0.0366	0.0143	0.0937	0.240	0.0586	0.1500	0.0229
PA1_rep1	0.0366	0.0143	0.0937	0.258	0.0586	0.2580	0.0229
PA1_rep2	0.0366	0.0143	0.0937	0.240	0.0586	0.1500	0.0229
PA1_rep3	0.0366	0.0143	0.1950	0.195	0.0762	0.0762	0.0229

	RF	SRE
PA1	0.3839	0
PA1_rep1	0.2580	0
PA1_rep2	0.3839	0
PA1_rep3	0.3839	0

```

$Sp290$PCA.median
      model.selected
PA1      "GBM"
PA1_rep1 "MARS"
PA1_rep2 "RF"
PA1_rep3 "RF"

$Sp290$thresholds
      PA1 PA1_rep1 PA1_rep2 PA1_rep3
prob.mean      684.7   651.5   643.5   579.5
prob.mean.weighted 592.3   610.7   602.1   505.2
median         695.1   650.8   568.3   511.8
Roc.mean       500.0   500.0   500.0   500.0
Kappa.mean     500.0   500.0   500.0   500.0
TSS.mean       500.0   500.0   500.0   500.0

```

OUTPUTS

Objects produced : `consensus.Future1.results` (in R's memory) which is the list returned by the function. It contains all the computational information that has been used to render the ensemble forecasts, the weights awarded to the models in the weighting process. The model selected by the PCA.median method (if set to True) is also returned and give us the model selected as the first axis of a PCA analyses (that means the model that explain the best the consensus probabilities). The forecasts themselves are stored on the hard disk directly in the corresponding folder.

NOTE1 : For the slot containing the weights (e.g. `$Sp281$weights`), the PA1 line corresponding to a run calibrate with all the pseudo-absences selected and presences data (models are evaluated on the same data so are often over optimistic). The PA1_rep1, PA1_rep2, and PA1_rep3 lines are linked to models calibrated and validated on two different subset of the pseudo-absences selected and presences data (DataSplit option in Models).

NOTE2 : The thresholds slot contains some consensus thresholds for different runs. `prob.mean` and `prob.mean.weighted` correspond respectively to the mean and weighted mean of thresholds used to convert probabilities into presences/absences data (e.g. (Evaluation.results.xx) table). `median` is the median of the same thresholds. The values of `Roc.mean`, `Kappa.mean` and `TSS.mean` is always set to 500. We made the assumption that as index are rescaled on a 0-1 ladder, 0.5 is the threshold that will discriminate presences and absences.

The function produces an object per species. These objects are arrays of three dimensions :

```
R code
load("proj.Future1/consensus_Sp290_Future1")
dim(consensus_Sp290_Future1)
```

```
R code
[1] 2264    4    6
```

```
R code
dimnames(consensus_Sp290_Future1)[-1]
```

```
R code
[[1]]
[1] "PA1"      "PA1_rep1" "PA1_rep2" "PA1_rep3"

[[2]]
[1] "prob.mean"      "prob.mean.weighted"
[3] "median"         "Roc.mean"
[5] "Kappa.mean"     "TSS.mean"
```

The second dimension is the repetition runs and the third dimension is the consensus methods. There is also an object called "Total_consensus_Future1" that makes a single output out of all the repetitions.

```
R code
load("proj.Future1/Total_consensus_Future1")
dim(Total_consensus_Future1)
```

```
R code
[1] 2264    2    6
```

```
R code
dimnames(Total_consensus_Future1)[-1]
```

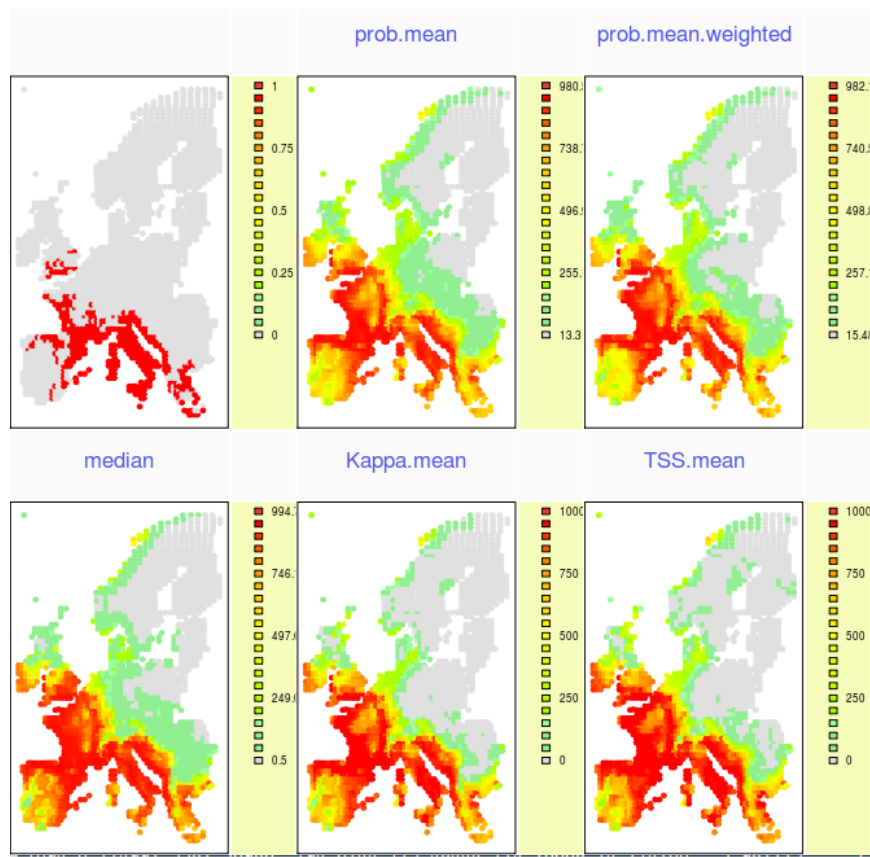
```
R code
[[1]]
[1] "Sp281" "Sp290"

[[2]]
[1] "prob.mean"      "prob.mean.weighted"
[3] "median"         "Roc.mean"
[5] "Kappa.mean"     "TSS.mean"
```

Now the second dimension is the species. Let's see and plot some of these :

R code

```
Total_consensus_Future1[1:20,,1]
multiple.plot(cbind(DataBIOMOD[, 'Sp281'],
                    Total_consensus_Future1[, 'Sp281',
                    c(1,2,3,5,6)]), LatLong, cex=0.8)
```



If binary is set to True, the same names are used with a terminal *_Bin* containing the consensus results in binary format.

7 Distributions Changes

7.1 Species Range Change

This function allows to estimate the proportion and relative number of pixels (or habitat) lost, gained and stable for the time slice considered : the range change.

The future range changes are calculated as a percentage of the species' present state. For example, if a species currently occupies 100 cells and is estimated by a model to cover 120 cells in the future, the range change will be + 20%.

The function uses two datasets. The current species distributions and the future one. Note that predictions for current and future must be in a binary (presence and absence) format and in the same resolution.

Let's use our data :

```
R code
load("proj.Future1/Total_consensus_Future1_Bin")
Biomod.RangeSize(CurrentPred = Sp.Env[,c(11,13)],
                  FutureProj = Total_consensus_Future1_Bin[, ,2],
                  SpChange.Save="SpChange")
```

A list of two datasets is created: Compt.By.Species and Diff.By.Pixel

Diff.By.Pixel stores useful information for each species. The species are in columns and the pixel in rows. For each species, a pixel could have four different values:

-2 if the given pixel is predicted to be lost by the species.

-1 if the given pixel is predicted to be stable for the species.

0 is the given pixel was not occupied, and will not be into the future.

1 if the given pixel was not occupied, and is predicted to be into the future.

In our examples :


```
R code
SpChange$Diff.By.Pixel[740:760,]
```

```
R code
      Sp281 Sp277
740     -1    -1
741      1     1
```

742	1	-1
743	1	1
744	0	-1
745	1	-1
746	0	-1
747	1	-1
748	1	-1
749	1	-1
750	0	-1
751	0	-1
752	-1	0
753	-1	0
754	-1	-2
755	-1	0
756	-1	-1
757	-1	-1
758	-1	-1
759	-1	-1
760	1	1

R code

```
multiple.plot(SpChange$Diff.By.Pixel, LatLong)
```



Compt.By.Species stores the summary of range change for each species (by rows).

The first four columns are relative numbers: Disa represents the number of pixels predicted to be lost by the given species. Stable0 is the number of pixels which are not currently occupied by the given species and not predicted to be. Stable1 represents the number of pixels currently occupied by the given species, and predicted to remain occupied into the future. Gain represent the number of pixels which are currently not occupied by the given species but predicted to be into the future.

PercLoss, PercGain and SpeciesRangeChange are the related percentage estimating as the following:

- CurrentRangeSize represent the modelled current range size (number of pixels occupied) of the given species.
- FutureRangeSize0Disp represents the future modelled range size assuming no migration of the given species.
- FutureRangeSize1Disp represents the future modelled range size assuming migration of the given species (depending on the datasets given in input, if Migration has been used or not).

SpChange\$Compt.By.Species *R code*

	<i>Loss</i>	<i>Stable0</i>	<i>Stable1</i>	<i>Gain</i>	<i>PercLoss</i>	<i>PercGain</i>
<i>Sp281</i>	9	1326	383	546	2.296	139.3
<i>Sp277</i>	94	550	986	634	8.704	58.7

	<i>SpeciesRangeChange</i>	<i>CurrentRangeSize</i>
<i>Sp281</i>	137	392
<i>Sp277</i>	50	1080

	<i>FutureRangeSize.NoDisp</i>	<i>FutureRangeSize.FullDisp</i>
<i>Sp281</i>	383	929
<i>Sp277</i>	986	1620

For other examples, we need some extra species data than the one we have been modelling. Load the dataset called DATA100SP.txt :

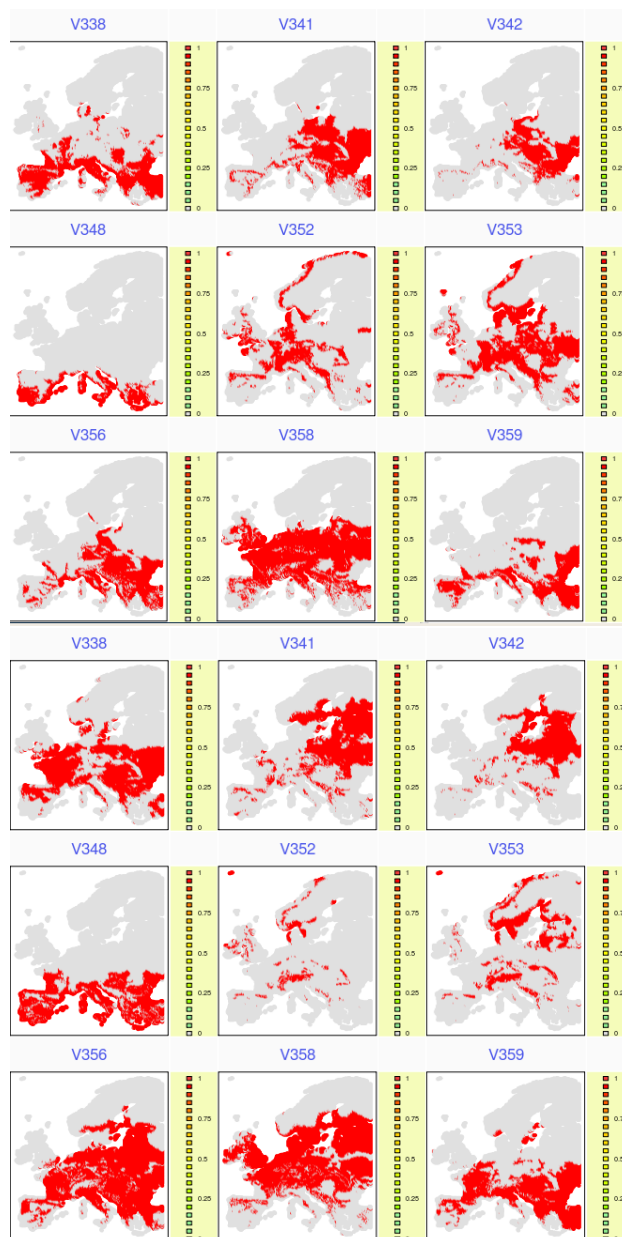
load("DATA100SP.RData")
#adds three objects : storeC, storeF and Curr
ls()

[1] "BestModelByRoc"
[2] "BestModelByTSS"
[3] "biomodDependencies"
[4] "Biomod.material"
[5] "Biomod.PA.data"
[6] "Biomod.PA.sample"
[7] "consensus_Sp290_Future1"
[8] "Curr"
[9] "DataBIOMOD"
[10] "DataEvalBIOMOD"
[11] "data.used"
[12] "Evaluation.results.Kappa"
[13] "Evaluation.results.Roc"
[14] "Evaluation.results.TSS"
[15] "Expl.Var"
[16] "Expl.Var2"
[17] "Expl.Var3"
[18] "Future1"
[19] "GBM.list"
[20] "GBM.perf"
[21] "i"
[22] "isnullYweights"
[23] "LatLong"
[24] "missingPackages"

```
[25] "model"
[26] "myPackages"
[27] "obj"
[28] "our.lines"
[29] "Pred"
[30] "Pred2"
[31] "Pred3"
[32] "PredBestModelByKappa"
[33] "Pred_Sp281"
[34] "Pred_Sp290"
[35] "Pred_Sp290_BinKappa"
[36] "Pred_Sp290_FiltKappa"
[37] "Pred_Sp290_indpdt"
[38] "Proj_Future1_Sp290"
[39] "Proj_Future1_Sp290_BinRoc"
[40] "rand"
[41] "Resp.Var"
[42] "Sp290_GLM_PA1"
[43] "Sp290_RF_PA1"
[44] "SpChange"
[45] "Sp.Env"
[46] "store"
[47] "storeC"
[48] "storeF"
[49] "Total_consensus_Future1"
[50] "Total_consensus_Future1_Bin"
[51] "VarImportance"
```

It corresponds to the run of the FDA on 100 species with the same resolution on current and future data, and coordinates (in Curr). Let's have a look at them :

```
multiple.plot(storeC[,1:9], Curr[,1:2], cex=0.7)
multiple.plot(storeF[,1:9], Curr[,1:2], cex=0.7)
```



Let's have a look at the SRC for some of those species :

R code

```

Biomod.RangeSize(CurrentPred = storeC,
                  FutureProj = storeF,
                  SpChange.Save="SpChange100")
SpChange100$Compt.By.Species[1:20,]

```

R code

	Loss	Stable0	Stable1	Gain	PercLoss	PercGain
V338	2568	16028	4347	6588	37.14	95.2711
V341	3988	16826	3215	5502	55.37	76.3848
V342	3710	19587	994	5240	78.87	111.3946
V348	306	22084	1840	5301	14.26	247.0177
V352	4426	23197	1269	639	77.72	11.2204
V353	7175	17449	2637	2270	73.12	23.1349
V356	809	13362	6030	9330	11.83	136.4235
V358	3980	11136	9629	4786	29.25	35.1679
V359	951	19655	4556	4369	17.27	79.3354
V360	3641	11938	10053	3899	26.59	28.4723
V365	606	9956	11092	7877	5.18	67.3363
V366	2752	23198	1151	2430	70.51	62.2598
V367	1016	16503	3358	8654	23.23	197.8509
V368	1347	11846	6820	9518	16.49	116.5422
V372	1053	21293	4553	2632	18.78	46.9497
V376	556	15217	4958	8800	10.08	159.5938
V377	591	23803	1602	3535	26.95	161.1947
V379	1817	27081	628	5	74.31	0.2045
V382	396	24532	1750	2853	18.45	132.9450
V385	400	18976	3045	7110	11.61	206.3861
	SpeciesRangeChange		CurrentRangeSize			
V338			58.134			
V341			21.019			
V342			32.526			
V348			232.759			
V352			-66.497			
V353			-49.990			
V356			124.594			
V358			5.923			
V359			62.066			
V360			1.884			
V365			62.156			
V366			-8.250			
V367			174.623			
V368			100.049			
V372			28.166			
V376			149.510			
V377			134.245			
V379			-74.110			
V382			114.492			
V385			194.775			
	FutureRangeSize.NoDisp		FutureRangeSize.FullDisp			
V338			4347			
V341			3215			
V342			994			
V348			1840			
V352			1269			
V353			2637			

V356	6030	15360
V358	9629	14415
V359	4556	8925
V360	10053	13952
V365	11092	18969
V366	1151	3581
V367	3358	12012
V368	6820	16338
V372	4553	7185
V376	4958	13758
V377	1602	5137
V379	628	633
V382	1750	4603
V385	3045	10155

R code

```
samp <- sample(100, 1)
x11()
level.plot(SpChange100$Diff.By.Pixel[,samp], Curr[,1:2], cex=0.6,
           title=colnames(storeC)[samp])
```

7.2 Species Turnover

This function allows to estimate species loss, gained, and turnover by pixel for the time slice considered.

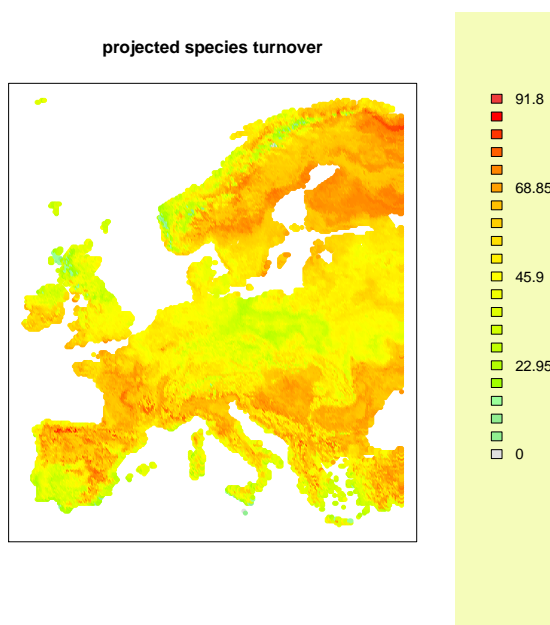
The function uses two datasets: the current species distributions and a future one. Note that predictions for current and future must be in a binary (presence and absence) format.

We can calculate the projected turnover for the 100 species and produce a plot of the turnover values.

```

R code
Biomod.Turnover(CurrentPred = storeC, FutureProj = storeF,
                 Turnover.Save= "Turnover")
level.plot(Turnover[,7], Curr[,1:2],
           title='projected species turnover',
           cex=0.6)

```



```

R code
Turnover[740:750,]

```

	Loss	Stable0	Stable1	Gain	PercLoss	PercGain	Turnover
740	15	59	18	8	45.45	24.24	56.10
741	15	58	19	8	44.12	23.53	54.76

742	16	52	21	11	43.24	29.73	56.25
743	10	61	23	6	30.30	18.18	41.03
744	9	61	24	6	27.27	18.18	38.46
745	11	58	25	6	30.56	16.67	40.48
746	9	59	26	6	25.71	17.14	36.59
747	9	59	27	5	25.00	13.89	34.15
748	10	57	26	7	27.78	19.44	39.53
749	8	60	26	6	23.53	17.65	35.00
750	10	58	25	7	28.57	20.00	40.48
<hr/>							
	CurrentSR	FutureSR.NoDisp	FutureSR.FullDisp				
740	33		18			26	
741	34		19			27	
742	37		21			32	
743	33		23			29	
744	33		24			30	
745	36		25			31	
746	35		26			32	
747	36		27			32	
748	36		26			33	
749	34		26			32	
750	35		25			32	

In the stored database, 10 columns are created.

The first four columns are relative numbers: Disa represents the number of species predicted to disappear from the given pixel. Stable0 is the number of species which are currently not in the given pixel and not predicted to migrate. Stable1 represents the number of species currently occurring in the given pixel, and predicted to remains into the future. Gain represent the number of species which are currently absent but predicted to migrate in the given pixel.

PercLoss, PercGain and Turnover are the related percentage estimated as the following:

- $\text{PercLoss} = 100 \times L / (\text{SR})$
- $\text{PercGain} = 100 \times G / (\text{SR})$
- $\text{Turnover} = 100 \times (L+G) / (\text{SR}+G)$

Where SR is the current species richness.

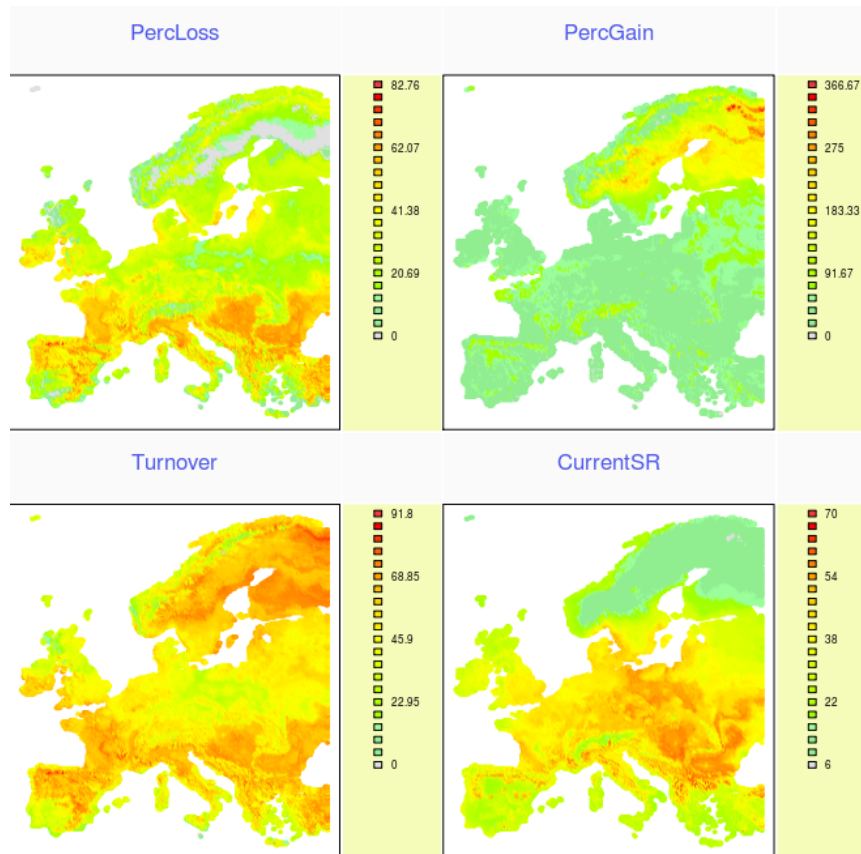
CurrentSR represent the current modelled species richness in the given pixel.

FutureSR0Disp represents the future modelled species richness assuming no migration of species

FutureSR1Disp represents the future modelled species richness assuming migration (depending on the datasets given in input, if Migration has been used or not).

R code

```
multiple.plot(Turnover[,5:8], Curr[,1:2], cex=0.6)
```



7.3 Probability Density Function

This function enables an overall viewing of the future projections range per species and gives the likelihood of range shift estimations. The optimal way for condensing 50, 75, 90 and 95% of the data will be calculated.

initial: a vector in a binary format (ones and zeros) representing the current distribution of a species which will be used as a reference for the range change calculations.

projection: a matrix grouping all the projections where each column is a single prediction. Make sure you keep projections in the same order as the initial vector (line1=site1, line2=site2, etc.).

Resolution: the step used for classes of projection in graphics. The default value is 5.

NOTE: modifying the resolution will directly influence the probability scale. Bigger classes will cumulate a greater number of predictions and therefore represent a greater fraction of the total predictions. The probability is in fact that of the class and not of isolated events.

cvsm: stands for current vs new. If true, the range change calculations will be of two types: the percentage of cells currently occupied by the species to be lost, and the relative percentage of cells currently unoccupied but projected to be, namely 'new' cells, compared to current surface range.

With the example above where the species will have 120 suitable sites in the future whilst only 100 at present, this might be the result of different events. A case could be that the 100 present cells are kept and an additional 20 new sites makes the 120 cells. Another possibility is that the 100 current cells are predicted to be lost with 120 new cells, also giving 120 total cells in future.

These two cases bring the same SRC calculations results, but whilst the first case does not imply much as in survival strategies (the current populations will still be in good conditions in future, plus even having new potential territories to explore and colonise), the second case, however, implies a strong migrating effort for the populations to stay in suitable environments. Those two cases and all in-between possibilities are distinguishable with this method.

groups: an option for ungrouping the projections enabling a separated visualisation of the prediction range per given group. A matrix is expected

where each column is a single projection and each line is giving details of one parameter.

Do keep in mind that this matrix represents the projections the way you have put them into the *projection* argument. Sort your matrix the way you have sorted your projections!

It can look like this:

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]
[1,]	"GAM"	"GAM"	"GAM"	"CTA"	"CTA"	"CTA"	"RF"	"RF"
[2,]	"Roc"	"Kappa"	"TSS"	"Roc"	"Kappa"	"TSS"	"Roc"	"Kappa"
	[,9]							
[1,]	"RF"							
[2,]	"TSS"							

```

#preparation of data
## Sp281
scenarios <- "Future1"
models <- c("ANN","CTA","GAM","GBM","GLM","MARS","FDA","RF","SRE")
evaluation <- c("Roc", "Kappa", "TSS")
reps <- 4
PAs <- c("PA1", "PA2")
DataFrame <- matrix(NA, 2264, 2)
Groups <- matrix(NA, 3, 216)
Groups[1,] <- c(rep(models,24))
Groups[2,] <- c(rep(rep(evaluation, each=36),2 ))
Groups[3,] <- c(rep(rep(PAs, each=108), 1))
for(sc in scenarios){
  for(PA in PAs){
    for(ev in evaluation){

      eval(parse(text=paste("load('proj.", sc, "/Proj_", sc,
                           "_Sp281_Bin", ev, "'")", sep="")))
      add.data <- eval(parse(text=paste("Proj_", sc, "_Sp281_Bin",
                                       ev, sep="")))

      DataFrame <- cbind(DataFrame, add.data[, , 'total.data', PA])
      DataFrame <- cbind(DataFrame, add.data[, , 'rep1', PA])
      DataFrame <- cbind(DataFrame, add.data[, , 'rep2', PA])
      DataFrame <- cbind(DataFrame, add.data[, , 'rep3', PA])
    }
  }
}

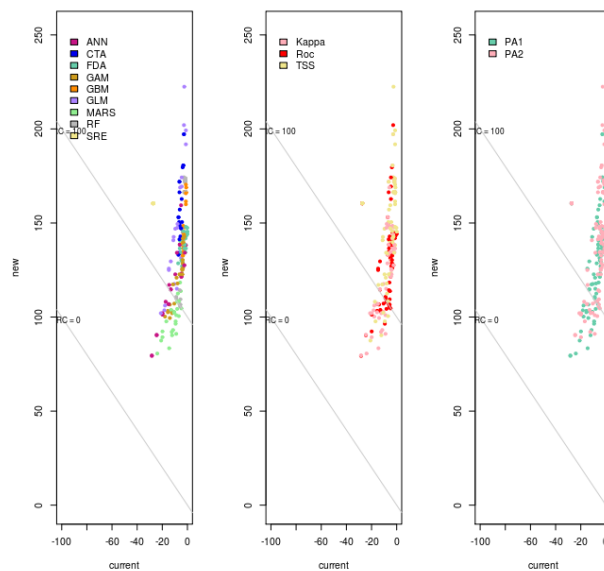
```

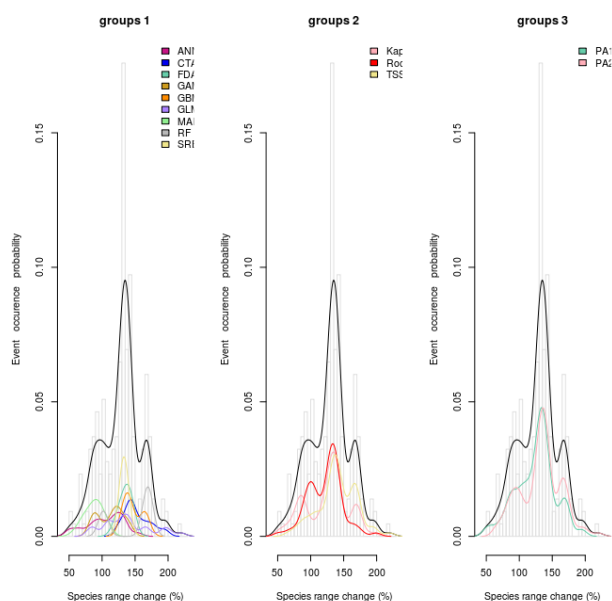
```
PDFdata281 <- DataFrame[, -1][, -1]
ProbDensFunc(initial=Sp.Env[, 'Sp281'], projections=PDFdata281, cvsn = T,
              groups = Groups, resolution = 5)
```

R code


\$stats

	<i>lower limit</i>	<i>upper limit</i>
50%	116.33	146.2
75%	80.87	147.2
90%	80.87	172.2
95%	66.07	177.3





Probability density function



The two lines represent where the SRC value is 0 (no absolute change in the number of suitable sites) and +100% (the species will double its current potential distribution size). Along those line, you have all the possibilities for giving that one value ($-10+10=0$; $-40+40=0$; ...).

On the cvsn graph, each dot is a projection. See how the single SRC value does not reflect every thing that is going on. In certain cases it hides the potential loss of current habitats, which would surely lead to different

management decisions if known.

7.3.1 An example with repetitions

The help file of the ProbDensFunc function provides a full example. It is done with 20 repetitions for half of the models to assess the variability in prediction making when the calibration of the model is done on partial data. Only Sp163 is done. Please look in details the help file for an example of the data preparation you should go through to run the function properly.

R code

`example(ProbDensFunc)`
